



Cyber-Physical Systems  
in Mechanical Engineering TU Berlin

# Applied Machine Learning in Engineering

**Lecture 01 summer term 2025**

Prof. Merten Stender

Cyber-Physical Systems in Mechanical Engineering, Technische Universität Berlin

[www.tu.berlin/cpsme](http://www.tu.berlin/cpsme)

[merten.stender@tu-berlin.de](mailto:merten.stender@tu-berlin.de)

# Organizational Details



Cyber-Physical Systems  
in Mechanical Engineering TU Berlin

- **Lecture:**
  - **Wednesday 10am-12pm CT, 90min, room EB 107**
  - In-person, no video recording available
- **Exercise:**
  - **Tuesday 12pm-02pm CT, 90min, room EB 202**
  - Hands-on coding exercises
  - Please bring your own computer
- **ISIS ([LINK](#)):**
  - Slides, lecture notes and exercise sheets will be published ahead in time
  - **Poll on active course participation:** your choice until April 30<sup>th</sup> (otherwise: **unsubscription**)
- **Exam:**
  - Written digital exam **at TU Berlin**. Dates: 21.07.2025 (08:00am); 09.10.2025 (10:00am)
  - Exam registration: **ONLY** through Moses (registration open!)

# Recap: Lecture 00



- **Uncertain model parameters**
  - Structural damping of metals, ...
  - Nonlinear behavior (elastomers stiffness, ...)
- **Inherent modeling assumptions and limitations**
  - Simplified constitutive models, ...
  - Idealized assumptions on homogeneity, ...
- **Speed and energy-efficiency**
  - Homogenization of heterogeneous materials
  - Low-order yet fast surrogate models

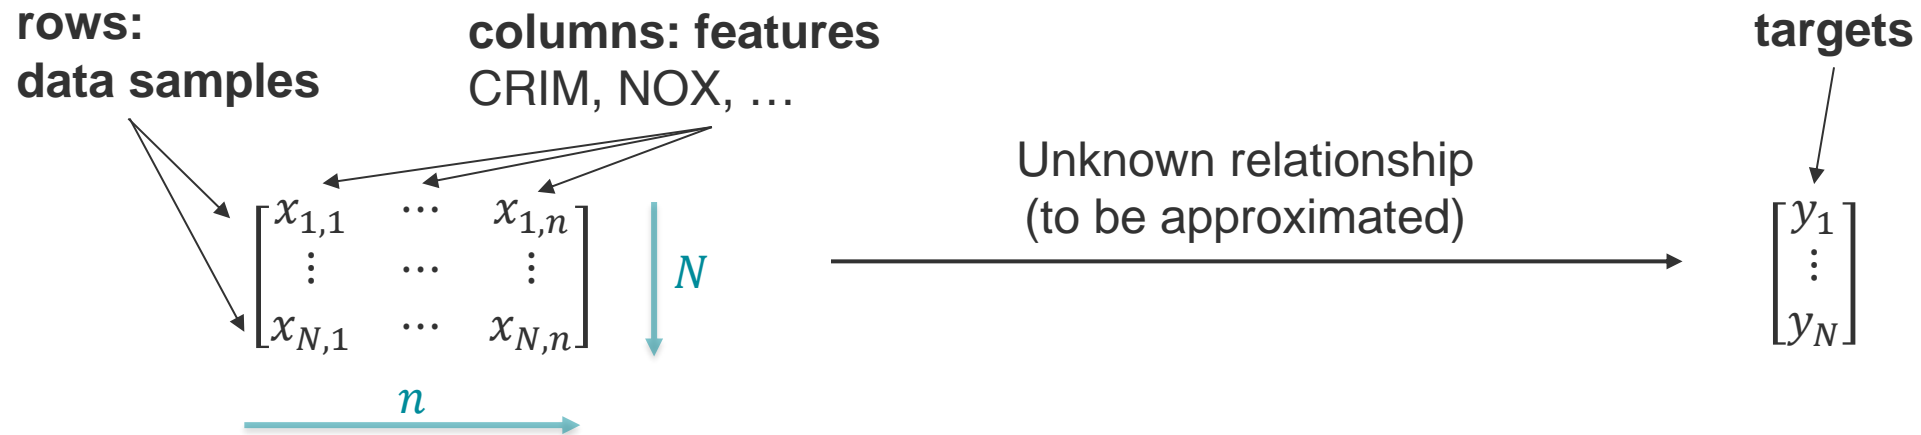
**Data-driven methods are particularly promising and powerful when a handcrafted algorithm is not existent or extremely difficult to formulate, parametrize, or execute.**

# Recap: Lecture 00



## Structured data (tabular data)

- **Features** (attributes): quantities that describe measurements or characteristic properties of an individual data sample (record)



- High-dimensional data sets:  $n$  very large
  - Big data:  $N$  very large (and  $n$  potentially, too)
- implies


# Recap: Lecture 00



## Unstructured data (also denoted as non-tabular data)

- Examples:
  - Text
  - Audio
  - Video
  - Images ...
- Special about unstructured data:
  - Additional latent dimensions
  - **Order matters** (latent dim.)

audio can be stored in an array but is not structured!



[ I | live | in | Munich | but | work | in | Berlin ]





Order of features not interchangeable!

# Recap: Lecture 00




## supervised learning (predictive task)




model training = reduce prediction error

## data (tabular)



## unsupervised learning (descriptive task)



finding clusters, groups and anomalies


# Recap: Exercise 00



- Set up a **Python 3.10+** environment for the semester
- Build some basic Python **programming skills** (following an [online tutorial](#)).
- Create a first **figure** using matplotlib

```
t = np.arange(start=0, stop=8 + 0.01, step=0.01) # time vector
freq = 10
x = np.pi * np.cos(freq * t) # x(t)

fig = plt.figure() # figsize=(8,4), dpi=200)
plt.plot(t, x, color='red', linewidth=2)
plt.xlabel(r'time $t$')
plt.ylabel(r'$x(t)$')
plt.title(fr'Signal $x(t) = \pi \cdot \cos(\{freq\}t)$')
plt.savefig('my_plot_of_cos_signal.png')
plt.show()
```



# Recap: Exercise 00



- Implement the basic **Newton scheme**

## Function $y=f(x)$

```
def f(x: float) -> float:
    return x ** 3 - 3 * x - 10

def dfdx(x: float) -> float:
    return 3 * x ** 2 - 3
```

## Ingredients (helper functions)

```
def newton_iter(xn: float, f, dfdx) -> float:
    return xn - (f(xn) / dfdx(xn))

def converged_f(xn: float, f) -> bool:
    return f(xn) < 10 ** (-7)

def converged_x(xn_1: float, xn: float) -> float:
    return np.abs(xn_1 - xn) < 10 ** (-4)
```

## Newton procedure

```
# initial guess of the zero
xn = 5
n = 0

# using the first convergence criterion on f(xn)
while not converged_f(xn, f) and n < 100:
    print(f'iteration {n}: xn={xn}')
    xn_plus1 = newton_iter(xn, f, dfdx)
    xn = xn_plus1
    n += 1

print(f'zero of f(x) is at x={xn}')
```



# Agenda



- Distance metrics for continuous attributes
- Linear regression fundamentals: Least squares method
- Evaluation of regression models

## Python

- Loops, functions
- Type hints
- PEP8 style guide
- Test-driven development
- scikit-learn library

# Learning outcomes



Cyber-Physical Systems  
in Mechanical Engineering TU Berlin

## Learn to ...

- Formulate a linear regression learning task
- Derive the closed-form solution to the least-squares linear regression
- Measure regression model prediction errors using the  $R^2$  metric
- Understand the limitations of metrics and the importance of visual evaluation
- Understand the principles of Test-Driven Development (TDD)

## Know about ...

- The Minkowski distance
- The meaning and interpretation of the coefficient of determination ( $R^2$ )
- Separating what and how code is doing during the implementation (TDD)



# Linear regression



- Continuous attributes allow for distance operations. Some distance metrics:

## Minkowski distance

$$d(\mathbf{x}, \mathbf{y}) = (\sum_{i=1}^n |\mathbf{x}_i - \mathbf{y}_i|^r)^{1/r}$$

- Euclidean distance ( $L_2$ )

$$\|\mathbf{q}\|_2 = (\sum_{i=1}^n |q_i|^2)^{1/2} \quad \text{for distance vector } \mathbf{q} = \mathbf{x} - \mathbf{y}, \mathbf{x}, \mathbf{y} \in \mathbb{R}^n$$

- Manhattan distance ( $L_1$ )

$$\|\mathbf{q}\|_1 = \sum_{i=1}^n |q_i|$$

- Supremum distance ( $L_\infty$ )

$$\|\mathbf{q}\|_\infty = \lim_{r \rightarrow \infty} (\sum_{i=1}^n |q_i|^r)^{1/r}$$

# Regression – Definition




**Definition:** Regression is the task of learning a target function  $f$  that maps each attribute set  $x$  into a continuous-valued output  $y$ .

[Tan, Introduction to Data Mining]

- **Goal:** find a target function  $f(x)$  with minimal error on training data
- **Error:** different definitions of error available

- **Methods:**
  - Linear regression
  - Decision Trees
  - Neural Networks
  - ...





# Linear Regression



- Find target function  $f$  that obtains a linear relationship between
  - (explanatory) variables  $x$
  - and target value  $y$
- Target function, i.e. lin. regression model, is parameterized in  $\theta = [\theta_0, \dots]^T$ :  $f(\theta)$
- Data set  $D = \{(x_i, y_i) \mid i = 1, \dots, N\}$ ,  $x_i = [x_{i,1}, x_{i,2}, \dots, x_{i,n}]^T \in \mathbb{R}^n$ ,  $N$  samples,  $n$ -dim. feature space


- Scalar example:  $x \in \mathbb{R}^1$



# Linear Regression



- Scalar variable  $x$  and scalar target value  $y \rightarrow$  find  $f(x) = \theta_0 + \theta_1 x = \mathbf{x}_i^\top \boldsymbol{\theta}$  with  $\mathbf{x} = [1, x_i]^\top$
- **Generating process**  $y_i = \theta_0 + \theta_1 x_i + \epsilon_i = \mathbf{x}_i^\top \boldsymbol{\theta} + \epsilon_i$  for  $i = 1, \dots, N$  samples
- Unobserved random variable  $\epsilon$  causing deviations from a perfectly linear relationship
- **Prediction:**  $\hat{y} = f(x)$  (target function evaluated at  $x$ )
- **Prediction error:**  $\mathcal{L} = \sum \|y_i - \hat{y}_i\| = \sum \|\mathbf{y}_i - \mathbf{x}_i^\top \boldsymbol{\theta}\|$
- **Error metrics:**  $\|\cdot\|$  : absolute error, squared error, ...
- Squared error (sum of squares)  $\rightarrow$  **Least squares linear regression**



# Least Squares Linear Regression



- Solving a linear regression task for the minimum **sum of squared errors (SSE)**

- Sum of squared errors:  $\mathcal{L}_{\text{SSE}} = \sum_i (y_i - \hat{y}_i)^2, i = 1, \dots, N$

- Least squares method:  $\min_{\theta} \|y_i - \mathbf{x}_i^T \theta\|$  for data samples  $D = \{(\mathbf{x}_i, y_i) \mid i = 1, \dots, N\}$ ,

- Optimal model parameters  $\theta^*$

- Loss for scalar setting:  $\mathcal{L}(\theta, D) = \sum_i (y_i - \hat{y}_i)^2 = \sum_i (y_i - \theta_0 - \theta_1 x_i)^2$

- Minimum of SSE: basic calculus. Vanishing gradient of  $\mathcal{L}$  with respect to model parameters  $\theta$

$$\frac{\partial \mathcal{L}}{\partial \theta_0} = 0 \text{ and } \frac{\partial \mathcal{L}}{\partial \theta_1} = 0$$



# Least Squares Linear Regression



- Loss for scalar setting:

$$\mathcal{L} = \sum_i (y_i - \hat{y}_i)^2 = \sum_i (y_i - \theta_0 - \theta_1 x_i)^2$$

- Vanishing gradient of  $\mathcal{L}$

$$\frac{\partial \mathcal{L}}{\partial \theta_0} = -2 \sum_i (y_i - \theta_0 - \theta_1 x_i) = 0$$

$$\frac{\partial \mathcal{L}}{\partial \theta_1} = -2 \sum_i (y_i - \theta_0 - \theta_1 x_i) x_i = 0$$


pen & paper exercise

- Normal equation

$$\underbrace{\begin{bmatrix} N & \sum_i x_i \\ \sum_i x_i & \sum_i x_i^2 \end{bmatrix}}_{A} \underbrace{\begin{bmatrix} \theta_0 \\ \theta_1 \end{bmatrix}}_{x} = \underbrace{\begin{bmatrix} \sum_i y_i \\ \sum_i y_i x_i \end{bmatrix}}_{b}$$

$Ax = b$

→ Solve for  $\theta$  to find optimal parameters  $\theta^*$



# Least Squares: Multi-Regression



- Feature space is multi-dimensional  $\mathbf{x} \in \mathbb{R}^n, n > 2, \mathbf{x} = [1, x_1, \dots, x_n]^\top, D = \{(\mathbf{x}_i, y_i), i = 1, \dots, N\}$

- Linear multi-regression:  $\hat{y}_i = \mathbf{x}_i^\top \boldsymbol{\theta}, \quad \mathbf{x}_i = [1, x_{i,1}, \dots, x_{i,n}]^\top, \mathbf{x}, \boldsymbol{\theta} \in \mathbb{R}^n$   
full data set:  $\hat{\mathbf{y}} = \mathbf{X}\boldsymbol{\theta}$

- Sum of squares  $\mathcal{L} = \|\mathbf{y} - \hat{\mathbf{y}}\|^2$

- Solution:  $\boldsymbol{\theta}^* = \underset{\boldsymbol{\theta}}{\operatorname{argmin}} \mathcal{L}(D, \boldsymbol{\theta})$

# Task: Compute Normal Form



- Loss\*  $\mathcal{L} = \|\hat{\mathbf{y}} - \mathbf{y}\|^2$
- Model predictions  $\hat{\mathbf{y}} = \mathbf{X}\boldsymbol{\theta}$
- Minimum:  $\frac{\partial \mathcal{L}}{\partial \boldsymbol{\theta}} = \mathbf{0}$

- Compute solution  $\boldsymbol{\theta}^*$

\*Inversed formulation of the L2 loss w.r.t. previous slides (without any effect on result)

$$\frac{\partial \mathcal{L}}{\partial \boldsymbol{\theta}} = 0$$

$$\frac{\partial}{\partial \boldsymbol{\theta}} \|\mathbf{X}\boldsymbol{\theta} - \mathbf{y}\|_2^2 = 0$$

$$\frac{\partial}{\partial \boldsymbol{\theta}} (\mathbf{X}\boldsymbol{\theta} - \mathbf{y})^\top (\mathbf{X}\boldsymbol{\theta} - \mathbf{y}) = 0$$

$$\frac{\partial}{\partial \boldsymbol{\theta}} (\boldsymbol{\theta}^\top \mathbf{X}^\top - \mathbf{y}^\top) (\mathbf{X}\boldsymbol{\theta} - \mathbf{y}) = 0$$

$$\frac{\partial}{\partial \boldsymbol{\theta}} (\boldsymbol{\theta}^\top \mathbf{X}^\top \mathbf{X}\boldsymbol{\theta} - \boldsymbol{\theta}^\top \mathbf{X}^\top \mathbf{y} - \mathbf{y}^\top \mathbf{X}\boldsymbol{\theta} + \mathbf{y}^\top \mathbf{y}) = 0$$

$$\frac{\partial}{\partial \boldsymbol{\theta}} (\boldsymbol{\theta}^\top \mathbf{X}^\top \mathbf{X}\boldsymbol{\theta} - 2\boldsymbol{\theta}^\top \mathbf{X}^\top \mathbf{y} + \mathbf{y}^\top \mathbf{y}) = 0$$

$$2\mathbf{X}^\top \mathbf{X}\boldsymbol{\theta} - 2\mathbf{X}^\top \mathbf{y} = 0$$

$$\boldsymbol{\theta} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$$




# Measuring Regression Performance


# Visual Regression Model Evaluation




- Whenever building a regression model, it needs to be evaluated (on a hold-out data set)
- A visual evaluation tool is the scatter plot of expected (ground truth) vs. predicted values



example 1



example 2




# Measuring Goodness of Fit: $R^2$



- Linear regression is a **supervised learning** approach
  - True target values are known  $\rightarrow$  ground truth targets  $y_i$
  - Predictions  $\hat{y}$  can be compared against ground truth for measuring the goodness of fit

- Coefficient of determination**  $R^2 = 1 - \frac{\sum_i (y_i - \hat{y}_i)^2}{\sum_i (y_i - \bar{y})^2}$ ,  $i = 1, \dots, N$ , with sample mean  $\bar{y} = \frac{1}{N} \sum_i y_i$





- Residual sum of squares  $\sum_i (y_i - \hat{y}_i)^2$
- Total sum of squares  $\sum_i (y_i - \bar{y})^2$
- Perfect fit:  $R^2 = 1.0$
- Baseline model  $f(x) = \bar{y}$   $R^2 = 0.0$
- „Worse models“  $R^2 < 0.0$



# Anscombe's quartet



- Proposed by Francis Anscombe in 1973  
Graphs in Statistical Analysis". American Statistician. 27 (1): 17–21
- Four data sets with (approx.) **same statistics**
  - Mean ( $x, y$ )
  - Variance ( $x, y$ )
  - Correlation ( $x, y$ )
  - Linear regressor  $f(x) = \theta_0 + \theta_1 x + \epsilon$
  - $R^2$  of linear regressor



Wikipedia

## Take-away message:

- Always be cautious with metrics
- Always check model prediction results visually



# Python





# Python Basics

Loops, Functions, Type Hints

# Loops



- **FOR loop:** A for loop lets you repeat a block of code for each item in a sequence (like a list, string, or range of numbers).

```
for item in sequence:  
    # do sth. with item
```

```
fruits = ["apple", "banana", "cherry"]  
  
for fruit in fruits:  
    print("I like", fruit)
```

```
# range(n) gives numbers  
from 0 to n-1  
  
for i in range(5):  
    print(i)
```

- **WHILE loop:** A while loop keeps repeating as long as a condition is True

```
while condition:  
    # do something
```

```
count = 0  
  
while count < 3:  
    print("Count is", count)  
    count += 1
```

```
# This will never stop!  
  
while True:  
    print("Looping forever!")
```

# Functions



- Function: A function is a named block of code that does something. You can call it whenever you need it
  - Define this task once, then reuse it whenever you want
  - Functions help avoid repetition and organize code

```
def function_name():  
    # code to run
```

```
def say_hello():  
    print("Hello!")
```

```
say_hello()  
say_hello()
```

```
>>Hello!  
>>Hello!
```

```
def greet(name):  
    print(f"Hello, {name}!")
```

```
greet('Bob')  
greet('Anna')
```

```
>>Hello, Bob!  
>>Hello, Anna!
```

```
def square(x):  
    return x * x
```

```
result = square(4)  
print(result)
```

```
>>16
```

- Can, but does not require, one/multiple return objects
- Naming convention: **snake case** my\_fancy\_function()

# Type Hints



- Recap Exercise 00: Plotting call wrapped into a function with type hints

```
"""
Extra: wrap it into a function, accepting amplitude and frequency and returning
the plot
"""

def plot_cos_signal(amplitude: float = 1.0, frequency: float = 1.0) -> None:
    t = np.arange(start=0, stop=8 + 0.01, step=0.01) # time vector
    x = amplitude * np.cos(frequency * t) # x(t)

    plt.figure(figsize=(8, 4), dpi=200)
    plt.plot(t, x, color='red', linewidth=2)
    plt.xlabel(r'time $t$')
    plt.ylabel(r'$x(t)$')
    plt.title(fr'Signal $x(t) = \{amplitude\} \cdot \cos(\{frequency\}t)$')
    plt.savefig('extra_plot_of_cos_signal.png')
    plt.show()

# use the function to plot a different cosine signal
plot_cos_signal(amplitude=0.5, frequency=3.14)
```

function definition

arguments

default values

type hints

call of the function



# PEP8 Style Guide

# PEP 8 – Style Guide for Python Code



Cyber-Physical Systems  
in Mechanical Engineering TU Berlin

## Coding conventions for Python code. Standard style guide helps avoiding common errors

- Improves collaboration with other developers (everyone talking in same accent)
- Link: <https://peps.python.org/pep-0008/> ; better to read presentation at <https://pep8.org/>
- PEP: Python Enhancement Proposals – style guide is constantly evolving over time!

## Preliminary guides

### **Never** use as a variable name:

- Capital or small-cap „o“ → can be mixed up with zero
- Capital „i“ → can be mixed up with 1
- Small-cap „L“ → can be mixed up with 1 or capital „i“

# PEP 8 Naming Conventions (Basics)



- **Python scripts**

- Module: lower-case words separated by underscore
- Package: lower-case words, no separators

```
my_module.py  
mypackage.py
```

- **Variables**

- Small-cap characters or words, underscore-separated
- CONSTANTS: all-capital letters

```
x = 5  
my_variable = 'ten'  
PI = 3.14
```

- **Functions**

- No single characters
- Small-cap words, underscore-separated

```
def my_function():  
    pass
```

- **Classes**

- Camel-case (capitalize each word, no separator)
- Methods: syntax like functions

```
class Car():  
class SportsCar():
```

# PEP 8 Syntax Conventions



- Whitespaces around operators, and around „=“
- Line indentation and line breaks
- ... and many more!
- Further reading: [Link](#)

```
# Correct:  
spam(ham[1], {eggs: 2})
```

```
# Wrong:  
spam( ham[ 1 ], { eggs: 2 } )
```

```
# Correct:  
if x == 4: print(x, y); x, y = y, x
```

```
# Wrong:  
if x == 4 : print(x , y) ; x , y = y , x
```

```
# Aligned with opening delimiter.  
foo = long_function_name(var_one, var_two,  
                           var_three, var_four)
```





# Test-Driven Development



## Why write code that tests other pieces of code?

### Ensure correct functionality

- Does a piece of code know what we request it to do?
- Does code catch edge cases and undesired usage?

### Continuous integration (version control, Git, ...)

- Automatically checking functionality after code changes
- Pushing to code base only when passing tests

### Documentation and collaboration

- Good tests can replace long documentation
- Splitting responsibilities: requirement definition and implementation

# Test-Driven Development (TDD)



Paradigm in software development (among others)

**Separate *what* the code needs to do from *how* it does it**

- Focus on
  - Modular software: interfaces much more important than the implementation
  - Perspective of a user that knows only the interface of a piece of code (arguments, returns)
  - Thinking of behavior rather than of specific implementation details
  - Higher software quality, robustness, and maintenance readiness
- 3-stage procedure: **RED – GREEN – REFACTOR**
  1. **Red:** implement tests and make sure that all of them are failed
  2. **Green:** implement functionalities until all tests are passed
  3. **Refactor:** clean up and optimize code

# Test-Driven Development (TDD)



Cyber-Physical Systems  
in Mechanical Engineering TU Berlin

- Five steps to follow:
  1. Write tests for each feature that you require
  2. Run the tests and make sure that all tests fail (**red**)
  3. Write the simplest code that passes the tests
  4. Make sure that all tests pass now (**green**)
  5. Refactor and improve the code from step 3 and repeat from step 4 (**refactor**)

Example: Implement a function that adds two numeric numbers (floats, real-valued) using TDD

Interface definition:

```
def adding_function(a, b):  
    return sum_a_b
```

Code snippet: → [live demo](#)

- test\_add\_fun\_unittest.py
- add\_fun.py

# Test-Driven-Development: Result



```
import unittest
from add_fun import adding_function as add

class TestAddingFunction(unittest.TestCase):

    def test_floats(self) -> None:
        # test the addition of two floats
        self.assertAlmostEqual(add(1.5, 1.5), 3)

    def test_ints(self) -> None:
        # test the addition of two ints
        self.assertEqual(add(1, 2), 3)

    def test_negatives(self) -> None:
        # test adding a negative and a positive number
        self.assertEqual(add(-1, 1), 0)

    def test_types(self) -> None:
        # test whether an exception is raised for non-
        # numeric inputs
        # using the context manager
        with self.assertRaises(TypeError):
            add(5, 'five')

if __name__ == "__main__":
    # use the main to run this script directly from your
    # editor
    unittest.main()
```

## Checks (using unittest package)

- Correct result for
  - floats
  - ints
  - positives and negatives
- Check edge cases / improper usage
  - ValueError raised for string inputs

```
def adding_function(a, b):
    # Add function

    # catch some errors if no numbers are handed over
    if (type(a) is not float) and (type(a) is not int):
        raise TypeError('input a is not of type int or float')

    if (type(b) is not float) and (type(b) is not int):
        raise TypeError('input b is not of type int or float')

    return a + b
```

# scikit-learn library



Cyber-Physical Systems  
in Mechanical Engineering TU Berlin

- **scikit-Learn** is the most relevant library for classic machine learning in Python
  - Data preprocessing
  - Wide range of ML models
  - Model evaluation
  - Well-documented with underlying mathematics and literature
- [R2 score](#): `sklearn.metrics.r2_score(y_true, y_pred)`

```
pip install scikit-learn
```

The image shows a banner from the scikit-learn website. On the left, there is the scikit-learn logo (a blue circle with an orange 's' and the word 'learn' in orange) followed by navigation links: 'Install', 'User Guide', 'API', 'Examples', 'Community' (with an external link icon), and 'More' (with a dropdown arrow). The main part of the banner has a blue background with the text 'scikit-learn' in large white letters, and 'Machine Learning in Python' in smaller white italicized letters below it. Two orange buttons are at the bottom left: 'Getting Started' and 'Release Highlights for 1.6'. On the right, an orange background contains a list of bullet points.

- Simple and efficient tools for predictive data analysis
- Accessible to everybody, and reusable in various contexts
- Built on NumPy, SciPy, and matplotlib
- Open source, commercially usable - BSD license



# Exercise 01



# Least-Squares Regression



1. Implement your own version of a scalar linear regression function using `numpy` and the normal form derived in class.
2. Estimate the effective rolling resistance factor of a car from measurements of vehicle speed and engine power
  - Underlying physics: force  $F_{\text{wind}} = c_W A \cdot \frac{\rho_{\text{air}} \cdot v_{\text{rel}}^2}{2}$ ,  $F_{\text{roll}} = c_R M g \cos \alpha$   
power  $P = v \cdot F$
  - Unknown random variables in the measurements (wind velocity, road inclination  $\alpha$ )
  - Compute the R2 value of your fit and validate with `scikit-learn`
3. Using test-driven development, implement the R2 error metric





# Questions?