



Lower Bounds:

Covering Maps & Neighborhood graphs





Port numbering model



You are in a room
with three doors,
labeled 1, 2, and 3.


>

—




You are in a room
with three doors,
labeled 1, 2, and 3.
> open door 3









You enter a room with two doors, labeled 1 and 2. You just came in through doorway 1.



You enter a room with
two doors, labeled 1
and 2. You just came
in through doorway 1.

>





High-level plan

- **Goal:**
 - show that problem X cannot be solved in the port-numbering model
- **General approach:**
 - construct port-numbered networks so that some nodes u, v, \dots will always produce the *same output*
 - show that if u, v, \dots have the same output, then it is *not a feasible solution* for X

High-level plan

- **Goal:**

- show that problem X cannot be solved in the port-numbering model





Covering
maps used
here

- **General approach:**

- construct port-numbered networks so that some nodes u, v, \dots will always produce the *same output*
- show that if u, v, \dots have the same output, then it is *not a feasible solution* for X

Covering map

- **Two port-numbered networks:**
 - $N = (V, P, \{p_v\}_{v \in V})$
 - $N' = (V', P', \{p'_v\}_{v \in V'})$
- **Surjection $f: V \rightarrow V'$ that preserves:**
 - inputs
 - degrees
 - connections
 - port numbers



Covering map

- “Fools” any deterministic algorithm
- If f is a covering map from N to N' , then:
 - v and $f(v)$ have the same state *before* round 1
 - v and $f(v)$ send the same messages in round 1
 - v and $f(v)$ receive the same messages in round 1
 - v and $f(v)$ have the same state *after* round 1


Covering map

- “Fools” any deterministic algorithm
- If f is a covering map from N to N' , then:
 - v and $f(v)$ have the same state *before* round T
 - v and $f(v)$ send the same messages in round T
 - v and $f(v)$ receive the same messages in round T
 - v and $f(v)$ have the same state *after* round T


Common steps

- Starting point: graph problem X
- Which graph G would be a “hard instance”?
- How to choose a port numbering N of G ?
- How to choose the other network N' ?
- How to construct mapping from N to N' ?


Example: 2-node path




Example: 4-node path



Example: cycles



Example: spanning trees



Common setup

- *N is the network we care about*
 - simple port-numbered network
 - well-defined and interesting underlying graph
- *N' is something strange*
 - not necessarily a simple port-numbered network
 - running A in N' makes no sense
 - introduced only to analyze what happens when we run A in N

LOCAL model

**What can you
do in T rounds in
the LOCAL model?**

Best that we can do?

- No restrictions on message size

Best that we can do?

- No restrictions on message size
- No restrictions on local computation


Best that we can do?

- No restrictions on message size
- No restrictions on local computation
- Possible to do:
*in each round, tell each neighbor
everything you know!*


Best that we can do?

- No restrictions on message size
- No restrictions on local computation
- Possible to do, and best that you can do:
*in each round, tell each neighbor
everything you know!*


At best:
in T rounds, each node can
learn its radius- T neighborhood



Round 1




Round 2




Round 3

Your **local output**
is a function of your
local neighborhood



local neighborhood

T -round algorithm is just a
mapping from radius- T neighborhoods
to local outputs



Running time =
number of communication rounds
until all nodes stop and produce
their local outputs


=

Locality =
how far do you need to see
in the graph to choose
your own part of the solution

New plan

- Algorithm A runs in T rounds and solves problem X
 - A is a mapping from radius- T neighborhoods to local outputs
- Construct examples where such a mapping cannot solve X correctly
 - Problem X is not solvable in T rounds

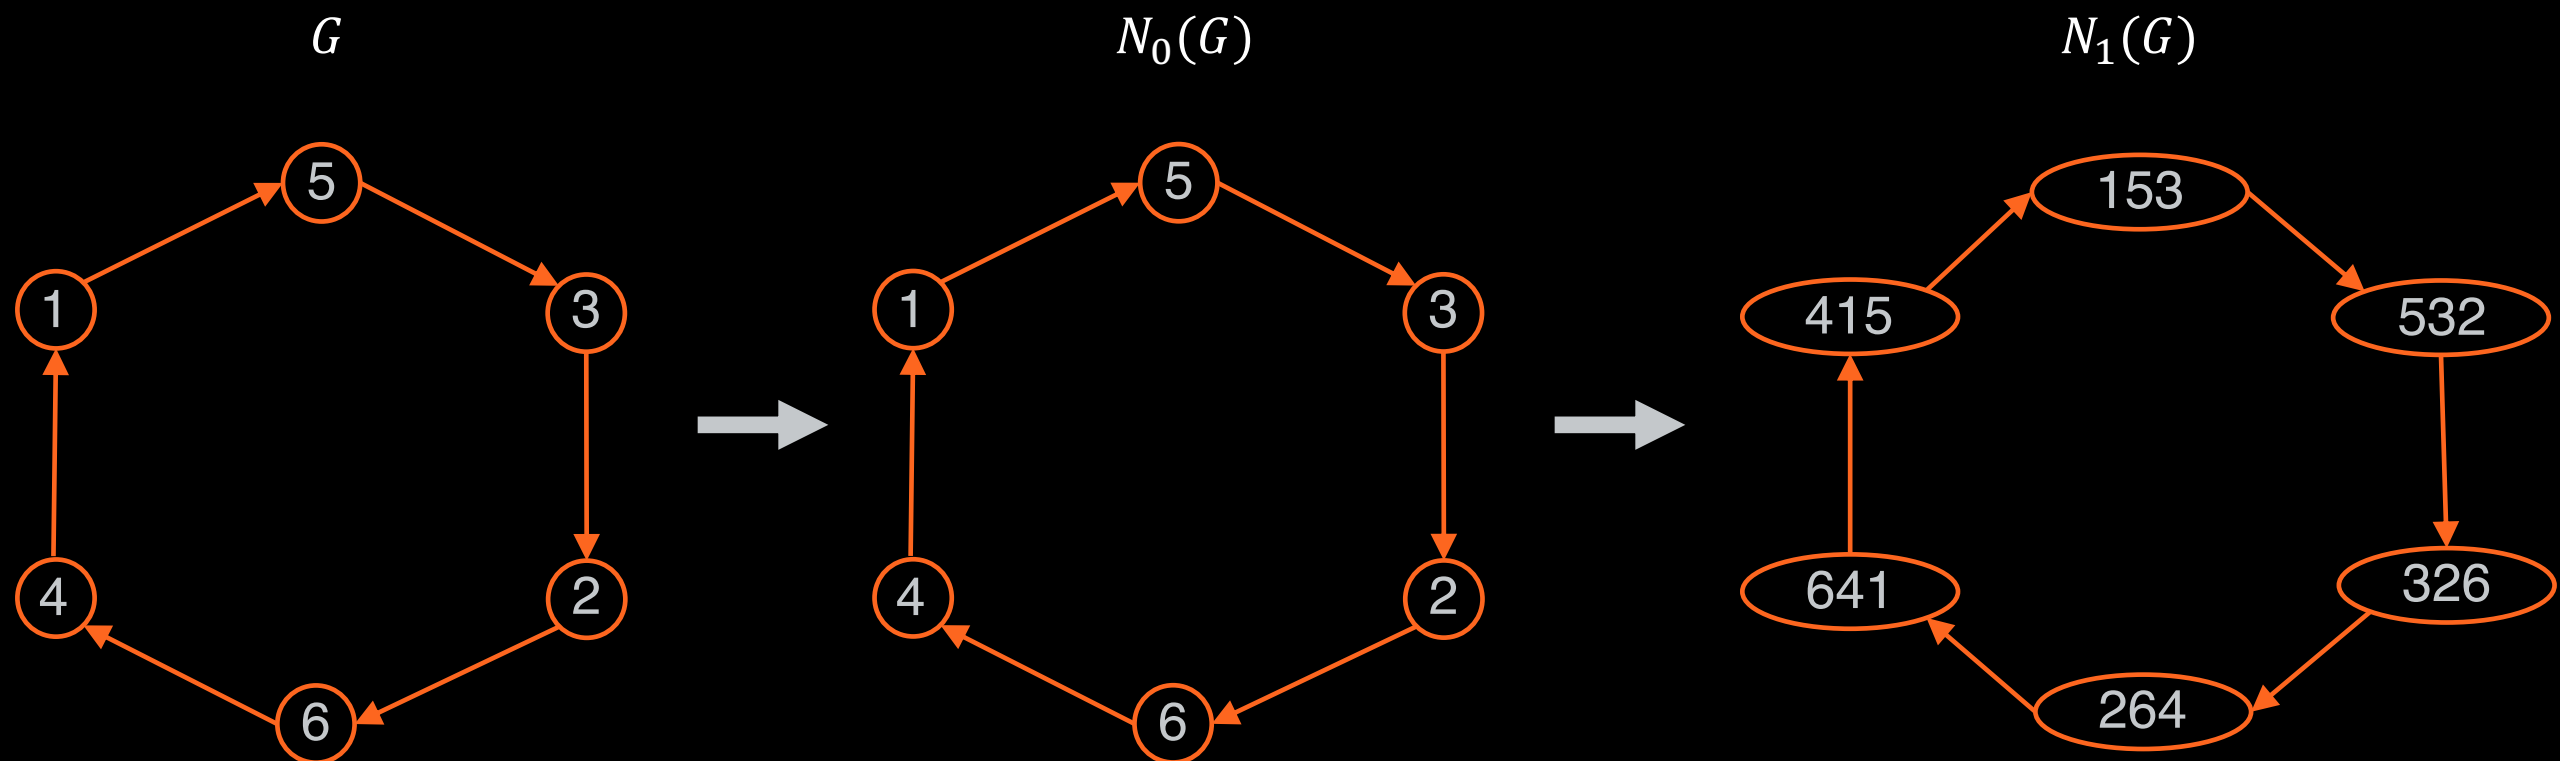
Example: 2-coloring a path



**What about 3-coloring a
~~path~~ directed cycle?**

Neighborhood graph

The r -neighborhood graph $N_r(G)$ consists of all r -hop views of G (for all nodes) which are connected iff they could originate from two adjacent nodes.






Neighborhood graph

Let \mathcal{G} be a graph class. The r -neighborhood graph $N_r(\mathcal{G})$ consists of all r -hop views of \mathcal{G} (for all nodes) which are connected iff they could originate from two adjacent nodes.

\mathcal{G} are cycles with nodes 1-6

$N_0(\mathcal{G})$?

$N_1(\mathcal{G})$?




These are just subgraphs of the respective neighborhood graphs!


Neighborhood graph

Let \mathcal{G} be a graph class. The r -neighborhood graph $N_r(\mathcal{G})$ consists of all r -hop views of \mathcal{G} (for all nodes) which are connected iff they could originate from two adjacent nodes.


\mathcal{G} are cycles with nodes 1-6



$N_0(\mathcal{G})$




Subgraph of $N_1(\mathcal{G})$




Neighborhood graph

There is an r -round algorithm that colors graphs \mathcal{G} with c colors iff the chromatic number of the neighborhood graph is $\chi(N_r(\mathcal{G})) \leq c$.


\mathcal{G} are cycles with nodes 1-6



$N_0(\mathcal{G})$



Subgraph of $N_1(\mathcal{G})$




Example: 4-cycle



Neighborhood graph

Intuitively:
the chromatic number declines for larger r , in
logarithmic factors “per hop”

Subgraph of $N_1(\mathcal{G})$



Subgraph of $N_2(\mathcal{G})$



Lower bound idea

- $\chi(N_r(\mathcal{G}))$ gives us a lower bound on the number of colors that any algorithm must use
- With each “hop”, the chromatic number of $N_r(\mathcal{G})$ goes down by at most a log factor
- After r hops, the lower bound on the chromatic number is

$$c > \chi(N_r(\mathcal{G})) \geq \log^{r-1} n > \log^* n$$

Lower bound idea

- $\chi(N_r(\mathcal{G}))$ gives us a lower bound on the number of colors that are needed to color $N_r(\mathcal{G})$. Here one needs to consider a subgraph of the neighborhood graph, see lecture notes
- With each “hop”, the chromatic number of $N_r(\mathcal{G})$ goes down by at most a log factor
- After r hops, the lower bound on the chromatic number is

$$c > \chi(N_r(\mathcal{G})) \geq \log^{r-1} n > \log^* n$$