



Shared Memory


Spectrum of Distributed (Computer) Systems




E.g., graphical processing units (GPUs) and specialized devices, in which large arrays of *simple processors* work in lock-step (“Gleichschritt”), PRAM, ...




Multi-threaded + multi-core servers/desktops with *shared memory for communication*.



Loosely-coupled *peer-to-peer* systems with message passing communication



The Shared Memory Model



Shared memory consists of registers.

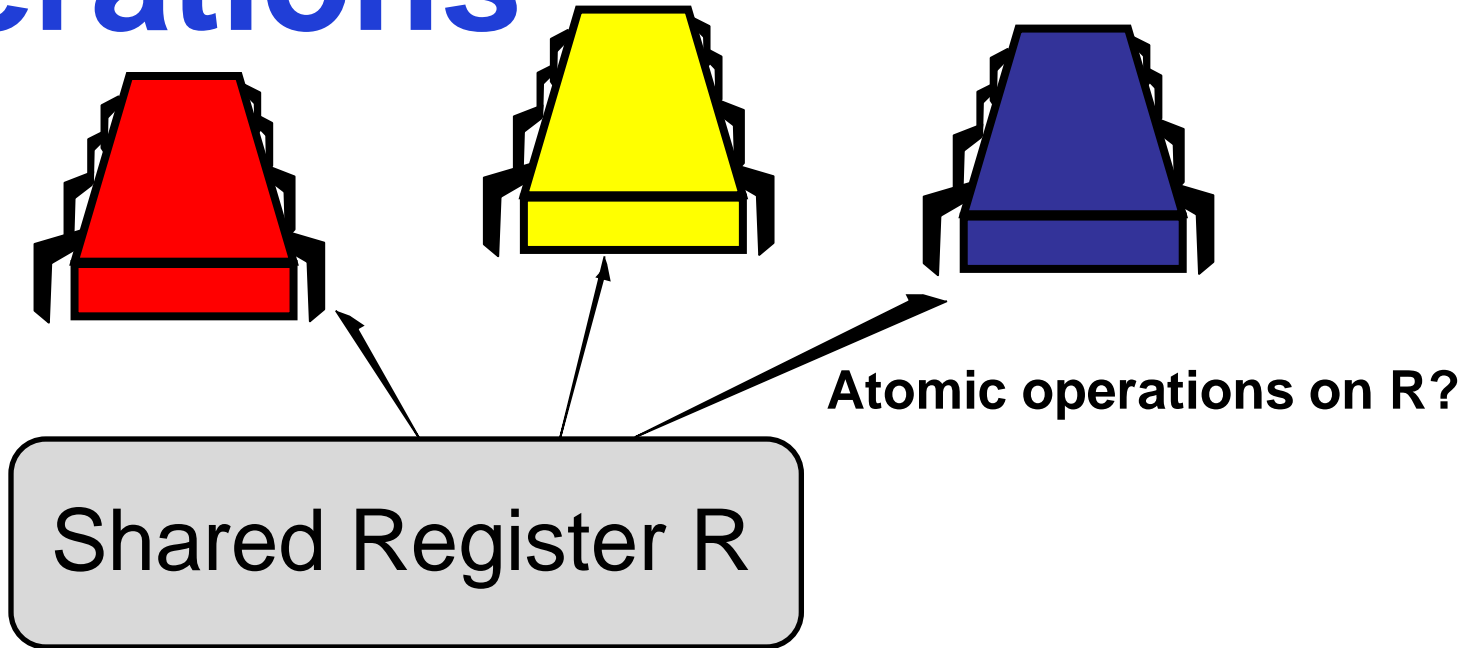
Formal Definition

Shared Memory

A shared memory system is a system that consists of asynchronous processes that access a *common (shared) memory*. A process can atomically access a register in the shared memory through a set of *predefined operations*. An atomic modification appears to the rest of the system instantaneously. Apart from this shared memory, processes can also have some *local (private) memory*.

Often a useful, simpler alternative model to reason about distributed systems!

Operations



Examples: (a.k.a. Data Types, Mealy Machines)

(1) **Test-and-Set(R):** $t := R$; $R := 1$; return t

(2) **Fetch-and-Add(R; x):** $t := R$; $R := R + x$; return t


(3) **Compare-and-Swap(R; x; y):**

if $R = x$ then $R := y$; return true;

else return false; endif

Why Shared Memory?

- Programming a shared memory system is *easier*: programmers access global variables directly!
- Because of this, even message passing systems often programmed through a *shared memory middleware*!
- From a message passing perspective, shared memory model is like a *bipartite graph*:



Nodes (asynchronous, may fail)

Shared registers (no failures, no delay)

The Power of RMW

The power of a shared memory system is determined by the *Consensus Number* (“universality of consensus”).

Consensus Number

The power of the RMW variant is measured by the consensus number. Consensus number k defines whether one can solve *consensus for k processes* (but not $k+1$).

Examples:

- *Test-and-Set* has consensus number 2
(one can solve consensus with 2 processes, but not 3)
- *Compare-and-Swap* has an infinite Consensus Numbers!

Desirable Properties of Distributed Systems

Safety, Liveness

Safety: “Something bad will never happen”,
Examples: some invariant holds
(function never returns -1 values),
serializability for DB transactions,
linearizability


Liveness: “Eventually something good happens”,
“system makes progress”

Precedes / Follows

Precedes / Follows

An operation op1 *precedes* an operation op2
iff op1 terminates before op2 starts.

An operation op2 *follows* operation op1 iff op1 precedes.



A Classic Shared Memory Problem



Fundamental *synchronization problem*: access to a resource

Mutual Exclusion

Each process executes the following code sections:

<Entry> → <Critical Section> → <Exit> → <Remaining Code>

A mutual exclusion algorithm consists of code for entry and exit sections, such that the following holds:

- (1) **Mutual Exclusion (Property?)**: At most one process is in the critical section.
- (2) **No deadlock (Property?)**: If some process manages to get to the entry section, later some (possibly different) process will get to the critical section.

Sometimes we in addition ask for

- (3) **No lockout (Property?)**: If some process manages to get to the entry section, later *the same process* will get to the critical section. (“Fairness”)
- (4) **Unobstructed exit (Property?)**: No process can get stuck in the exit section.

A Classic Shared Memory Problem



Fundamental *synchronization problem*: access to a resource

Mutual Exclusion

Each process executes the following code sections:

<Entry> → <Critical Section> → <Exit> → <Remaining Code>

A mutual exclusion algorithm consists of code for entry and exit sections, such that the following holds:

- (1) **Mutual Exclusion (Safety)**: At most one process is in the critical section.
- (2) **No deadlock (Liveness)**: If some process manages to get to the entry section, later some (possibly different) process will get to the critical section.

Sometimes we in addition ask for

- (3) **No lockout (Liveness)**: If some process manages to get to the entry section, later *the same process* will get to the critical section. (“Fairness”)
- (4) **Unobstructed exit (Liveness)**: No process can get stuck in the exit section.

Mutex

Test-and-Set(R): $t := R$; $R := 1$; return t

How to achieve Mutex with single Test-and-Set register?

Mutex

Test-and-Set(R): $t := R$; $R := 1$; return t

How to achieve Mutex with single Test-and-Set register?

Mutex

Input: Shared register $R:=0$

<Entry>

Repeat:

$r := \text{test-and-set}(R)$

Until $r=0$

<Critical Section>

...

<Exit>

$R:=0$

<Remaining Code>

...

} Set register to 1, then check whether it was so already.

Mutex

Test-and-Set(R): $t := R$; $R := 1$; return t

How to achieve Mutex with single Test-and-Set register?

Mutex

Input: Shared register $R := 0$

<Entry>

Repeat:

$r := \text{test-and-set}(R)$

Until $r = 0$

<Critical Section>

...

<Exit>

$R := 0$

<Remaining Code>

...

} Set register to 1, then check whether it was so already.

Correct Mutex?

No deadlock?

No lockout?

Unobstructed exit?

Mutex

- (1) Mutex: ?
- (2) Deadlock free: ?
- (3) Lockout: ?
- (4) Unobstructed exit: ?

Mutex

Input: Shared register $R:=0$

<Entry>

Repeat:

$r := \text{test-and-set}(R)$

Until $r=0$

<Critical Section>

...

<Exit>

$R:=0$

<Remaining Code>

...

Mutex

- (1) Mutex: **ok!**
- (2) Deadlock free: ?
- (3) Lockout: ?
- (4) Unobstructed exit: ?

Mutex

Input: Shared register $R:=0$

<Entry>

Repeat:

$r := \text{test-and-set}(R)$

Until $r=0$

<Critical Section>

...

<Exit>

$R:=0$

<Remaining Code>

...

Proof:

- (1) **Mutex:** R initially 0. Let p_i be the i -th process to successfully execute the test-and-set (i.e., result 0) at time t_i , and say at time t'_i , p_i resets $R:=0$. Between these times, nobody else can execute CS.

Mutex

- (1) Mutex: **ok!**
- (2) Deadlock free: **ok!**
- (3) Lockout: ?
- (4) Unobstructed exit: ?

Mutex

Input: Shared register $R:=0$

<Entry>

Repeat:

$r := \text{test-and-set}(R)$

Until $r=0$

<Critical Section>

...

<Exit>

$R:=0$

<Remaining Code>

...

Proof:

- (1) **Mutex:** R initially 0. Let p_i be the i -th process to successfully execute the test-and-set (i.e., result 0) at time t_i , and say at time t'_i , p_i resets $R:=0$. Between these times, nobody else can execute CS.
- (2) **Deadlock:** One of the processes waiting in the entry section will successfully test-and-set as soon as the process in the critical section exited.

Mutex

- (1) Mutex: **ok!**
- (2) Deadlock free: **ok!**
- (3) Lockout: ?
- (4) Unobstructed exit: **ok!**

Mutex

Input: Shared register $R:=0$

<Entry>

Repeat:

$r := \text{test-and-set}(R)$

Until $r=0$

<Critical Section>

...

<Exit>

$R:=0$

<Remaining Code>

...

Proof:

- (1) **Mutex:** R initially 0. Let p_i be the i -th process to successfully execute the test-and-set (i.e., result 0) at time t_i , and say at time t'_i , p_i resets $R:=0$. Between these times, nobody else can execute CS.
- (2) **Deadlock:** One of the processes waiting in the entry section will successfully test-and-set as soon as the process in the critical section exited.
- (4) **Exit:** Since the exit section only consists of a single instruction (no potential infinite loops) we have unobstructed exit.

Lockout

Test-and-Set(R): $t := R$; $R := 1$; return t

May be *unfair*!

Mutex

Input: Shared register $R:=0$

<Entry>

Repeat:

$r := \text{test-and-set}(R)$

Until $r=0$

<Critical Section>

...

<Exit>

$R:=0$

<Remaining Code>

...

} Always same process may win!
Solution?

Lockout

Test-and-Set(R): $t := R$; $R := 1$; return t

May be *unfair*!

Mutex

Input: Shared register $R:=0$

<Entry>

Repeat:

$r := \text{test-and-set}(R)$

Until $r=0$

<Critical Section>

...

<Exit>

$R:=0$

<Remaining Code>

...

Always same process may win!
Solution: make FIFO queue...

Lockout

Test-and-Set(R): $t := R$; $R := 1$; return t

May be *unfair*!

Mutex

Input: Shared register $R := 0$

<Entry>

Repeat:

$r := \text{test-and-set}(R)$

Until $r = 0$

<Critical Section>

...

<Exit>

$R := 0$

<Remaining Code>

...

Always same process may win!
Solution: make FIFO queue...

What about weaker objects?
Can I do without atomic RMW?

Lockout

Test-and-Set(R): $t := R$; $R := 1$; return t

May be unfair!

Mutex

Input: Shared register $R := 0$

<Entry>

Repeat:

$r := \text{test-and-set}(R)$

Until $r = 0$

<Critical Section>

...

<Exit>

$R := 0$

<Remaining Code>

...

Always same process may win!
Solution: make FIFO queue...


What about weaker objects?
Can I do without atomic RMW?

Yes: **Peterson's algorithm!**

Peterson's Algo

Assume: *two processes* only!

Need *three* registers (init: 0).



“P0 wants CS”

Written only by P0.

Read by both.

“P1 wants CS”


Written only by P1.

Read by both.

Peterson's Algo

Assume: *two processes* only!

Need *three* registers (init: 0).



“Who has priority at the moment?”

Written by both.

Peterson's Algo

Assume: *two processes* only!

Need *three* registers (init: 0).

Peterson's Mutex

Code for process P_i

<Entry>

$W_i := 1$

$\Pi := 1-i$

Loop until $\Pi = i$ or $W_{1-i} = 0$

<Critical Section>


...

<Exit>

$W_i := 0$

<Remaining Code>

...



Process indicates that it wants to enter CS in “*Want-Register*”. Can only do if other process does not want, or I have *priority* (*shared variable!*).

Peterson's Algo

Assume: *two processes* only!

Need *three* registers (init: 0).

Peterson's Mutex

Code for process P_i

<Entry>

$W_i := 1$

$\Pi := 1-i$

Loop until $\Pi = i$ or $W_{1-i} = 0$

<Critical Section>


...

<Exit>

$W_i := 0$

<Remaining Code>

...



Process indicates that it wants to enter CS in "*Want-Register*". Can only do if other process does not want, or I have *priority* (*shared variable!*).

Spin-Lock! (Busy-wait)

Priority register used to avoid deadlock!

Peterson

Peterson gives (1) mutex, (2) no deadlock, (3) no lockout, (4) unobstructed.

Proof:

- (1) **Mutex:** If both compete (“want”), only one can get priority and access CS.
- (2) **No Deadlock:** If both in loop and want, one process must have priority and it gets direct access to the critical section.
- (3) **Fairness:** Non-priority process waiting in loop gets priority when other process starts again! Shared variable.
- (4) **Unobstructed Exit:** Exit only a single instruction.

Peterson's Mutex

Code for process P_i

<Entry>

$W_i := 1$

$\Pi := 1-i$

Loop until $\Pi = i$ or $W_{1-i} = 0$

<Critical Section>

...

<Exit>


$W_i := 0$

<Remaining Code>

...

QED

Another fundamental task: Store&Collect



Goal: Collect up-to-date infos about other processes!
With *atomic RW* only.

Two operations:

- **sop(val)**: Process p_i *stores* val to be the latest value of its own register R_i . (“1:1 write”, so *single-writer model*)
- **cop()**: *Collects* a “view”, a function V where $V(p_i)$ is the latest value stored by p_i , for each process p_i .

Store&Collect

Assume: registers initialized to «?».

Note: Collect has no sequential specification and cannot be linearized.

Our goal here: A collect operation **cop** should never read from the future or miss a preceding store operation **sop**.

Store&Collect

For a collect operation **cop**, the following validity properties must hold for every process p_i :

1. If $V(p_i) = \text{"?"}$, then no store operation by p_i *precedes* **cop**.
2. If $V(p_i) = v$, then v is the value of a **sop** operation of p_i that does not *follow* **cop**, and there is no store operation by p_i that *follows* **sop** and *precedes* **cop**.

Complexity Measure

We measure the following complexity.

Step Complexity

Step complexity of an operation is the number of accesses to the registers in the shared memory.

How to implement a valid Collect() operation?

Simple Algorithm

Step **sop()** & **cop()**

Operation **STORE(val)**, by proc. p_i

$R_i := \text{val};$

Operation **COLLECT**:

for $i := 1$ to n do

$V(p_i) := R_i$ (* read register R_i *)

end

Works (atomic read/write).
Complexity?

Simple Algorithm

Step sop() & cop()

Operation **STORE(val)**, by proc. p_i

$R_i := \text{val};$

Operation **COLLECT**:

for $i := 1$ to n do

$V(p_i) := R_i$ (* read register R_i *)

end

Works (atomic read/write).

Complexity?

STORE is 1 step

COLLECT is n steps

Adaptive Algorithm

If only two processes wrote some value, COLLECT is too costly!
How to make an operation adaptive to the number of processes that were active in the execution?

Adaptive Operation

If up to time t , $k \leq n$ processes have started or finished at least one operation, an operation is called *adaptive* if step complexity depends on k but not on n .

How to make our algorithms adaptive?

Adaptive Algorithm

If only two processes wrote some value, COLLECT is too costly!
How to make an operation adaptive to the number of processes that were active in the execution?


Adaptive Operation

If up to time t , $k \leq n$ processes have started or finished at least one operation, an operation is called *adaptive* if step complexity depends on k but not on n .

How to make our algorithms adaptive?
We need *Splitters*...

Splitter

Splitter



Synchronization primitive:

- Process entering it exits with *stop, left* or *right*
- If k processes enter, at most one exists with *stop*, and *at most $k-1$* processes exit with *left* and at most $k-1$ processes exit with *right*.
- If single process enters it, *stop for sure*.

Not perfect balance, but there are two processes that obtain *different values* (stop, left, right).

How to implement splitter?

Splitter Algo

Splitter

Two shared registers $X: \{?, 1, \dots, n\}$, $Y: \text{bool}$

Initialization: $X=?, Y=\text{false}$

Splitter access by p_i :

$X := i$


If Y then *return right*

Else

$Y := \text{true}$

if $X=i$ then *return stop*

else *return left*



Splitter Algo

Splitter

Two shared registers $X: \{?, 1, \dots, n\}$, $Y: \text{bool}$
Initialization: $X=?, Y=\text{false}$

Splitter access by p_i :

$X := i$


If Y then *return right*

Else

$Y := \text{true}$

if $X=i$ then *return stop*

else *return left*



Why correct?

Correctness

A single process *always stops*:

- Clear: check solo-run

Splitter

Two shared registers $X: \{?, 1, \dots, n\}$, $Y: \text{bool}$

Initialization: $X=?, Y=\text{false}$

Splitter access by p_i :

$X:=i$

If Y then *return right*

Else

$Y:=\text{true}$

if $X=i$ then *return stop*

else *return left*

Correctness

A single process *always stops*:

- Clear: check solo-run

At most $k-1$ return right:

- **First process** checking Y will
not return right

Splitter

Two shared registers X: {?, 1, ..., n}, Y: bool

Initialization: X=?, Y=false

Splitter access by p_i :

X:=i

If Y then *return right*

Else

Y:=true

if X=i then *return stop*

else *return left*

Correctness

A single process *always stops*:

- Clear: check solo-run

At most $k-1$ return right:

- **First process** checking Y will *not return right*

At most $k-1$ return left:

- Assume process p is **last** to set $X:=i$
- If p does not return right, it will find its own value later and *stop*: it does *not return left*!

Splitter

Two shared registers $X: \{?, 1, \dots, n\}$, $Y: \text{bool}$

Initialization: $X=?, Y=\text{false}$

Splitter access by p_i :

$X:=i$

If Y then *return right*

Else

$Y:=\text{true}$

if $X=i$ then *return stop*

else *return left*

Correctness

A single process *always stops*:

- Clear: check solo-run

At most $k-1$ return right:

- **First process** checking Y will *not return right*

At most $k-1$ return left:

- Assume process p is **last** to set $X:=i$
- If p does not return right, it will find its own value later and *stop*: it does *not return left*!

At most one process stops:

- Assume contrary: both processes p_i and p_j return stop, and w.l.o.g. assume p_i sets $X:=i$ before p_j sets $X:=j$.
- Both can only reach “else” if Y was false for both! But then, X value of *p_i has been overwritten* in the meanwhile, and p_i does not return stop!

Splitter

Two shared registers $X: \{?, 1, \dots, n\}$, $Y: \text{bool}$

Initialization: $X=?, Y=\text{false}$

Splitter access by p_i :

$X:=i$

If Y then *return right*

Else

$Y:=\text{true}$

if $X=i$ then *return stop*

else *return left*

Correctness

Splitter

Two shared registers X: {?, 1, ..., n}, Y: bool

Initialization: X=?, Y=false

A single process *always stops*:

How to realize adaptive collect now?
Splitter trees!

else return left

At most k-1 return left:

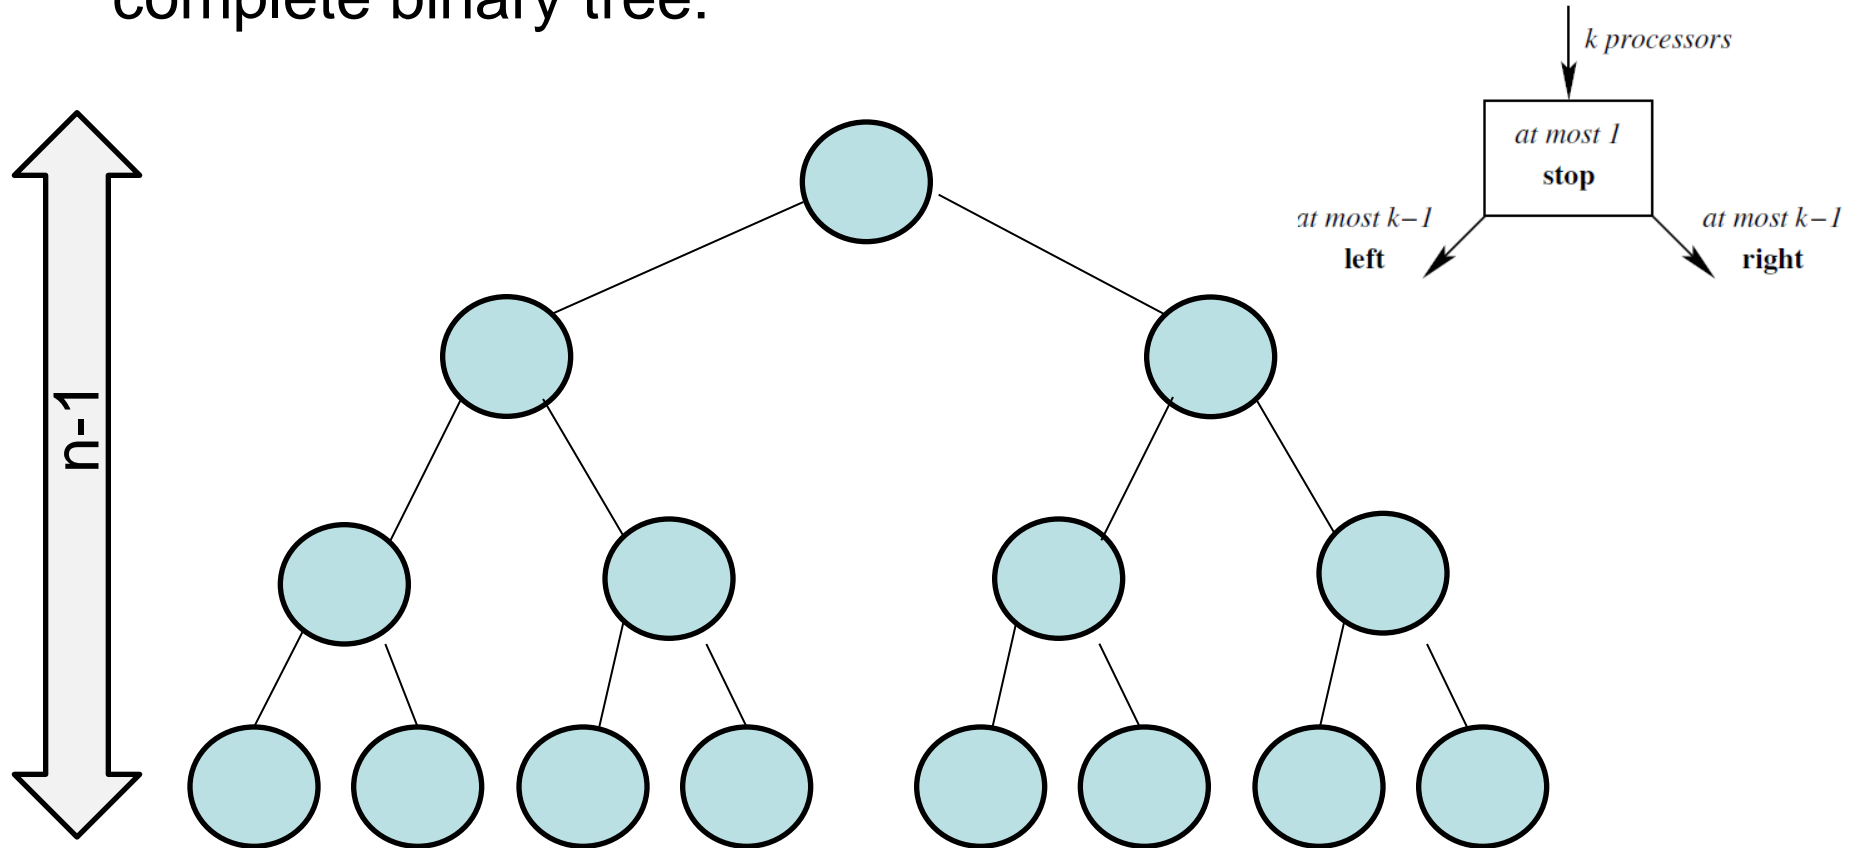
- Assume process p is **last** to set $X:=i$
- If p does not return right, it will find its own value later and *stop*: it does *not return left*!

At most one process stops:

- Assume contrary: both processes p_i and p_j return stop, and w.l.o.g. assume p_i sets $X:=i$ before p_j sets $X:=j$.
- Both can only reach “else” if Y was false for both! But then, X value of *p_i has been overwritten* in the meanwhile, and p_i does not return stop!

Splitter Tree

Assume we have $2^n - 1$ splitters, arranged in complete binary tree:



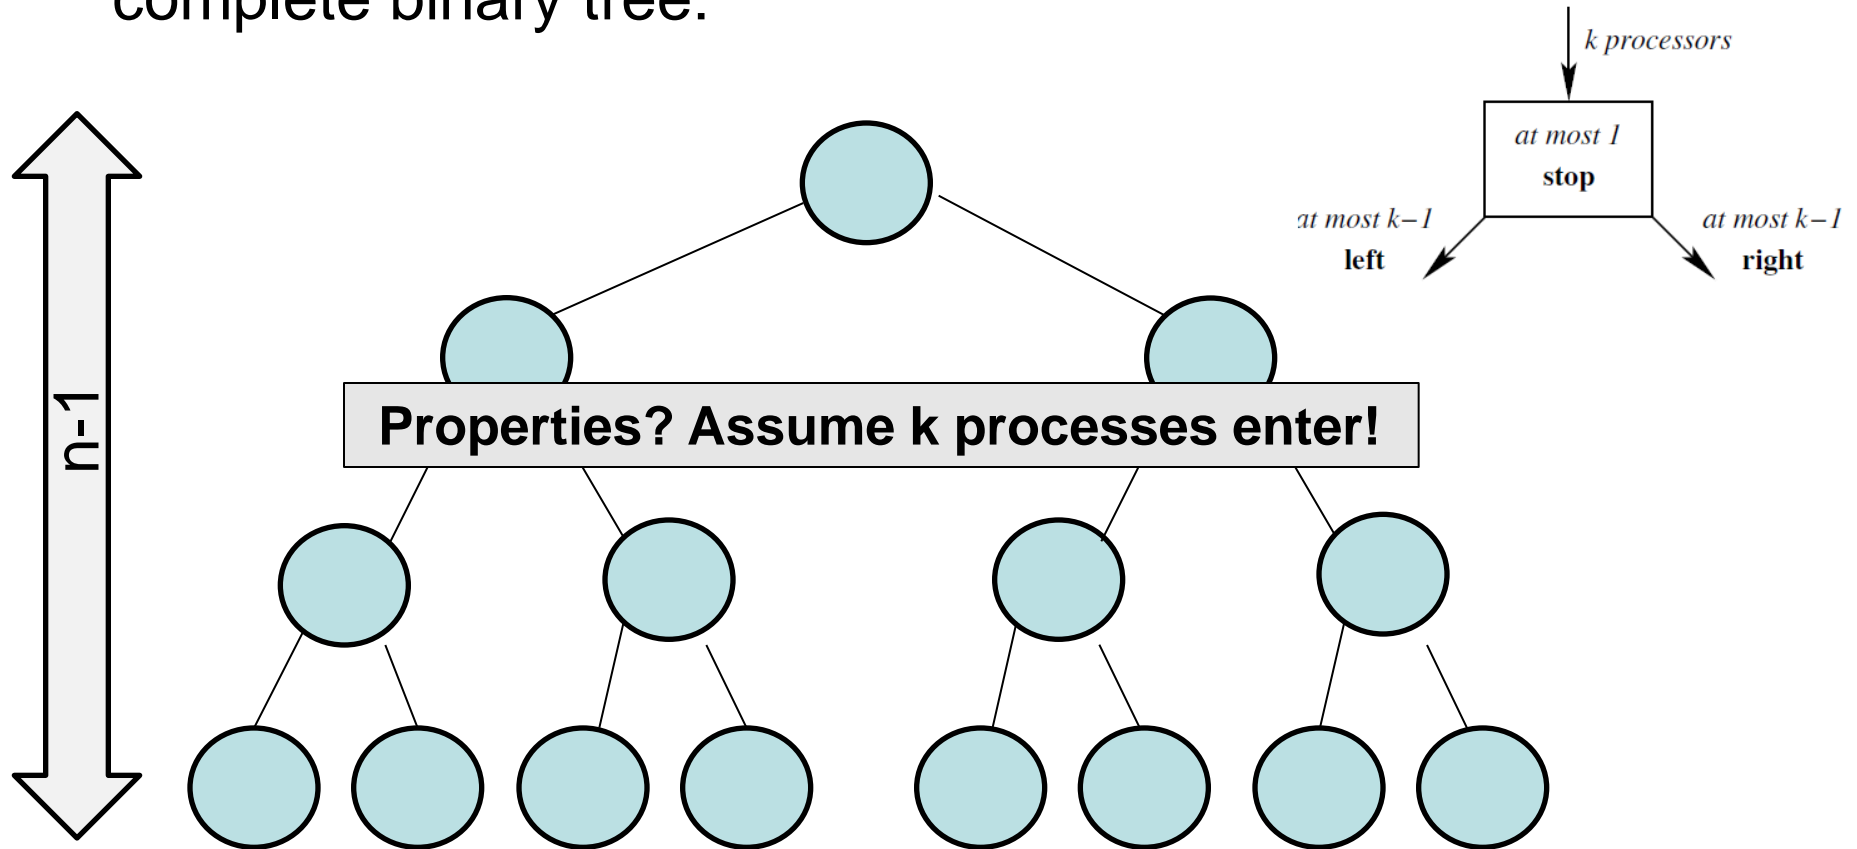
Let $S(v)$ be splitter of node v in tree.

Additionally: for every splitter, shared variables $Z_S: \{?, 1, \dots, n\}$ and boolean M_S .

A splitter is **marked** if $M_S = \text{true}$.

Splitter Tree


Assume we have $2^n - 1$ splitters, arranged in complete binary tree:



Let $S(v)$ be splitter of node v in tree.


Additionally: for every splitter, shared variables $Z_S: \{?, 1, \dots, n\}$ and boolean M_S .

A splitter is *marked* if $M_S = \text{true}$.



k processes *traversing splitter tree*:

- At most one process can stop at some given splitter
- Every process stops at some splitter at depth at most $k-1$. Why?

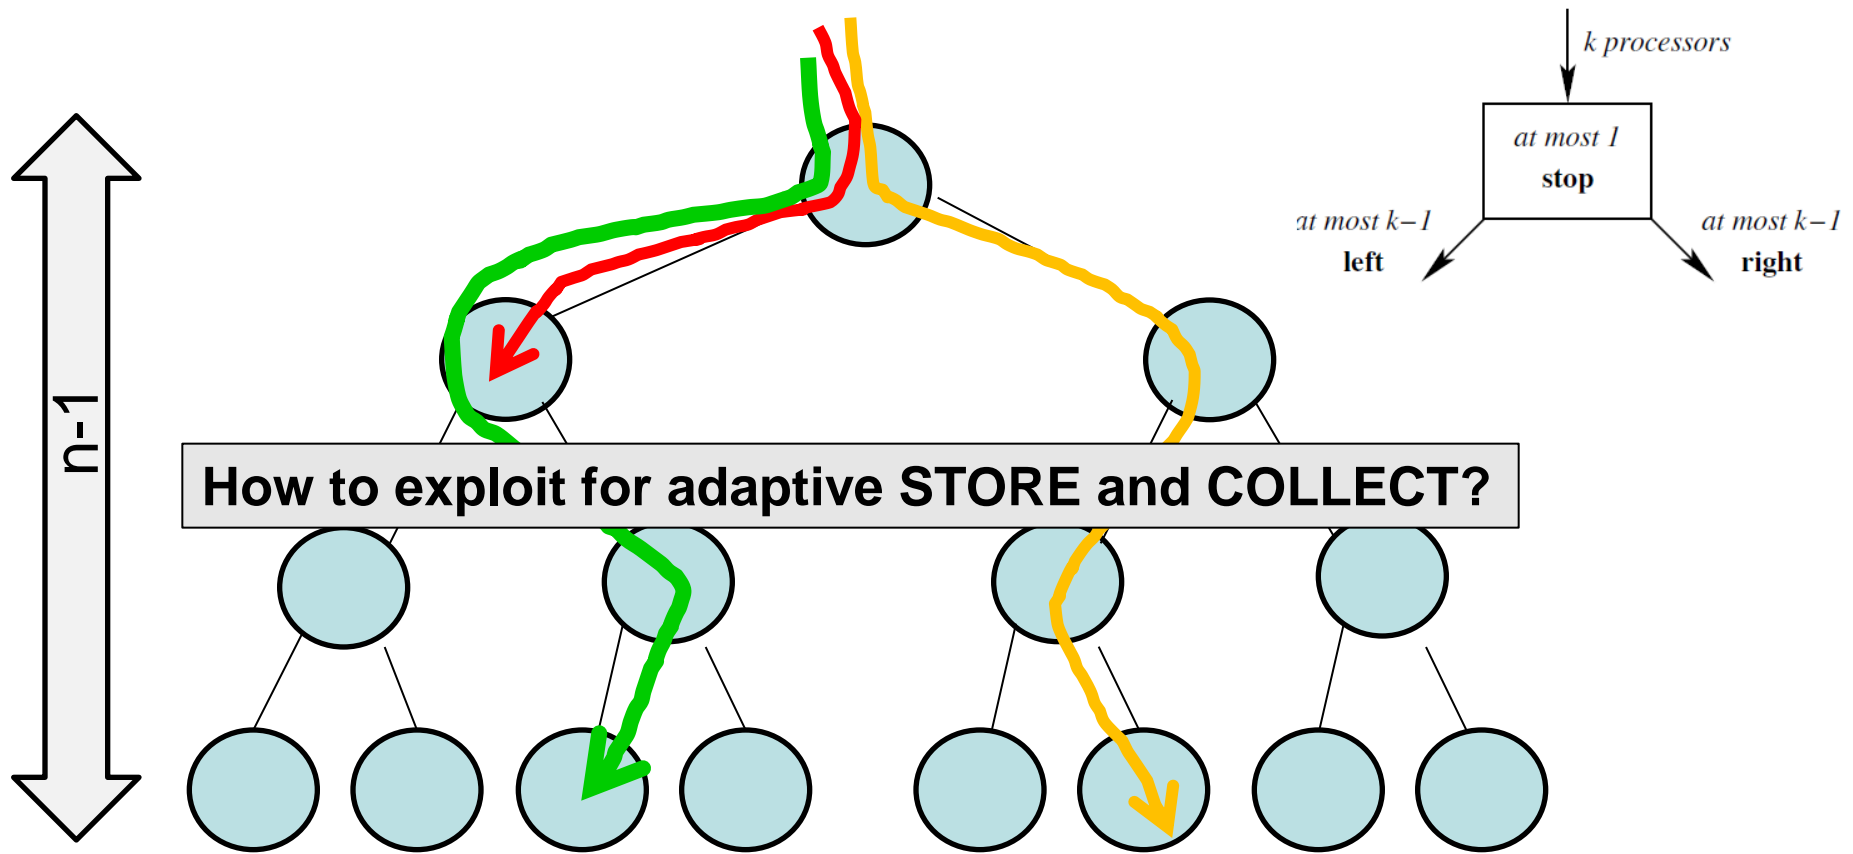


k processes *traversing splitter tree*:

- At most one process can stop at some given splitter
- Every process stops at some splitter at depth at most $k-1$. Why?

Proof by Induction:

- By definition, k processes enter the root splitter at depth 0
- If $k-i$ processes enter splitter at root of subtree at depth i (induction hypothesis), at most $k-i-1$ obtain left and at most $k-i-1$ right.



k processes *traversing splitter tree*:


- At most one process can stop at some given splitter
- Every process stops at some splitter at depth at most $k-1$. Why?

Proof by Induction:


- By definition, k processes enter the root splitter at depth 0
- If $k-i$ processes enter splitter at root of subtree at depth i (induction hypothesis), at most $k-i-1$ obtain left and at most $k-i-1$ right.

Upon *first store operation* of a process, *traverse* splitter tree:
data structure to mark which nodes were already active!

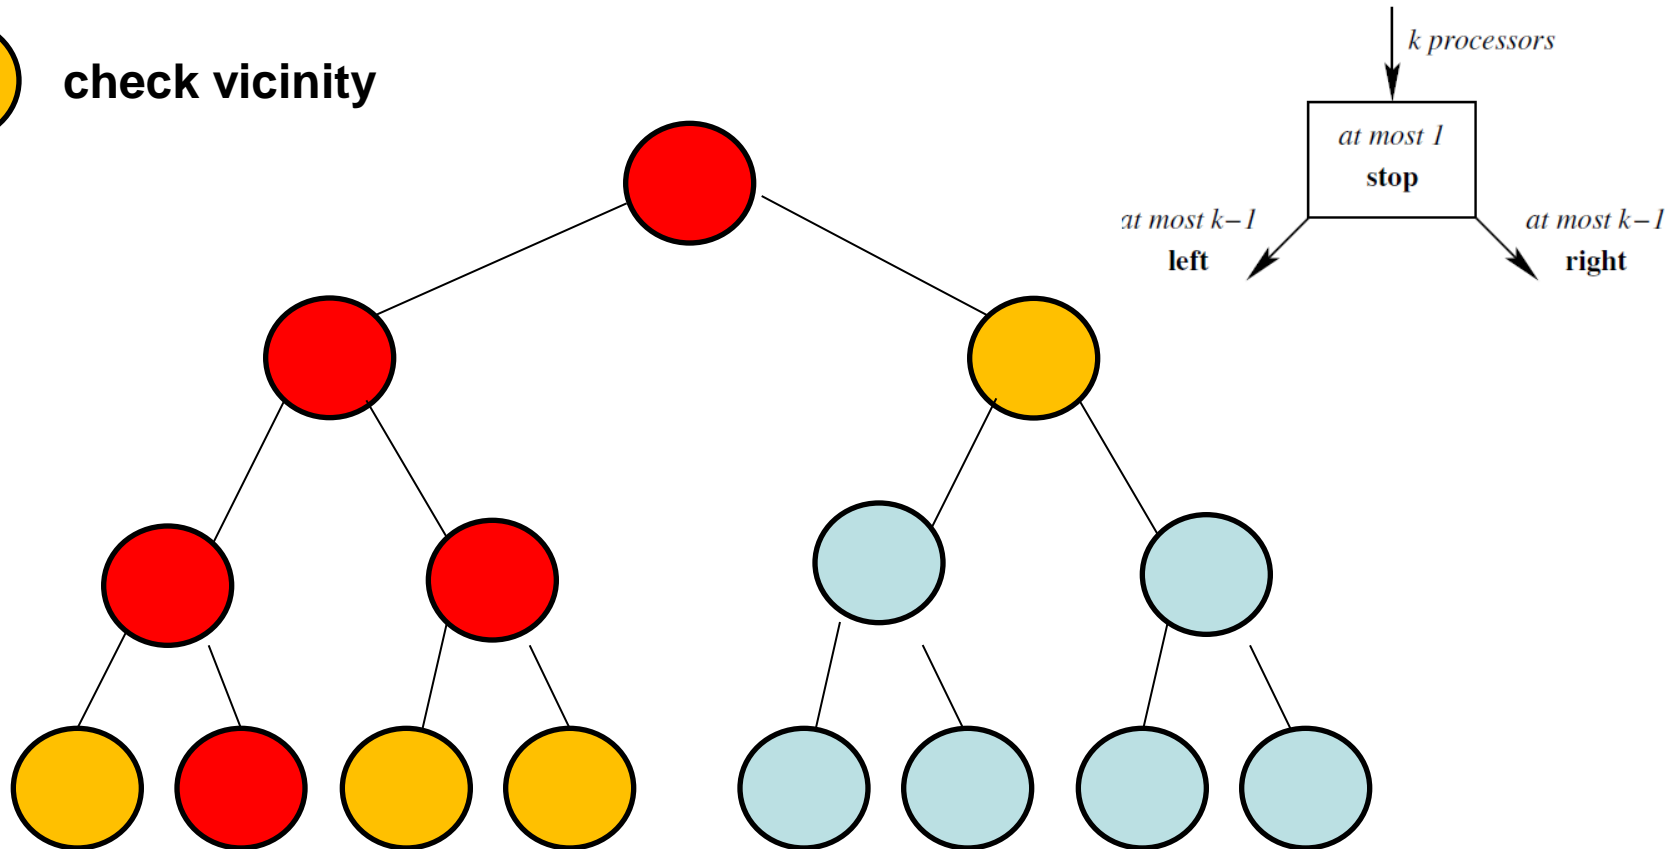
A connected component: *collect only needs to check this part!*



active



check vicinity



Tree Ops

Operation STORE(val) (by process p_i) :

```
1:  $R_i := val$ 
2: if first STORE operation by  $p_i$  then
3:    $v :=$  root node of binary tree
4:    $\alpha :=$  result of entering splitter  $S(v)$ ;
5:    $M_{S(v)} := \text{true}$ 
6:   while  $\alpha \neq \text{stop}$  do
7:     if  $\alpha = \text{left}$  then
8:        $v :=$  left child of  $v$ 
9:     else
10:       $v :=$  right child of  $v$ 
11:    end if
12:     $\alpha :=$  result of entering splitter  $S(v)$ ;
13:     $M_{S(v)} := \text{true}$ 
14:  end while
15:   $Z_{S(v)} := i$ 
16: end if
```

Traverse splitter tree from top, **mark traversed nodes (M)** and **store splitter** where p_i stopped.

Operation COLLECT:

Traverse marked part of binary tree:

```
17: for all marked splitters  $S$  do
18:   if  $Z_S \neq \perp$  then
19:      $i := Z_S$ ;  $V(p_i) := R_i$ 
20:   end if
21: end for
```

Only collect marked parts!

Tree Ops

Operation STORE(val) (by process p_i) :

```
1:  $R_i := val$ 
2: if first STORE operation by  $p_i$  then
3:    $v :=$  root node of binary tree
4:    $\alpha :=$  result of entering splitter  $S(v)$ ;
5:    $M_{S(v)} := \text{true}$ 
6:   while  $\alpha \neq \text{stop}$  do
7:     if  $\alpha = \text{left}$  then
8:        $v :=$  left child of  $v$ 
9:     else
```

Traverse splitter tree from top, mark traversed nodes (M) and store splitter

Complexity of solution?

```
11:   end if
12:    $\alpha :=$  result of entering splitter  $S(v)$ ;
13:    $M_{S(v)} := \text{true}$ 
14: end while
15:  $Z_{S(v)} := i$ 
16: end if
```

Operation COLLECT:

Traverse marked part of binary tree:

```
17: for all marked splitters  $S$  do
18:   if  $Z_S \neq \perp$  then
19:      $i := Z_S$ ;  $V(p_i) := R_i$ 
20:   end if
21: end for
```

Only collect marked parts!

Adaptive Collect

Step complexity of first STORE is $O(k)$, subsequent ones are $O(1)$. COLLECT has step complexity $O(k)$.

Proof.

- First *STORE*: splitter tree traversal to find “my” location, at most at depth k
- From then on, will always *STORE* there: $O(1)$
- *COLLECT*:
 - Only need to check marked part and their neighbors
 - Marked part of the tree is connected
 - At most $2k-1$ nodes are marked:
 - By induction: a single process entering a splitter will always stop
 - The right and left child of the root are subtrees too, first node will stop at first splitter.

Adaptive Collect

Step complexity of first STORE is $O(k)$, subsequent ones are $O(1)$. COLLECT has step complexity $O(k)$.

Proof.

- First **STORE**: splitter tree traversal to find “my” location, at most at depth k
- From then on, will always **STORE** there: $O(1)$
- **COLLECT**:
 - Only need to check marked part and their neighbors
 - Marked part of the tree is connected
 - At most $2k-1$ nodes are marked:
 - By induction: a single process entering a splitter will always stop
 - The right and left child of the root are subtrees too, first node will stop at first splitter.

Disadvantage?

Adaptive Collect

Step complexity of first STORE is $O(k)$, subsequent ones are $O(1)$. COLLECT has step complexity $O(k)$.


Proof.

- First **STORE**: splitter tree traversal to find “my” location, at most at depth k
- From then on, will always **STORE** there: $O(1)$
- **COLLECT**:
 - Only need to check marked part and their neighbors
 - Marked part of the tree is connected
 - At most $2k-1$ nodes are marked:
 - By induction: a single process entering a splitter will always stop
 - The right and left child of the root are subtrees too, first node will stop at first splitter.

Disadvantage? Space complexity! Store $O(2^n)$ tree in memory...

Splitter Matrix


Idea: Instead of arranging splitters in 2^n binary tree, arrange them in *$n \times n$ matrix*:



5x5 splitter matrix

Splitter Matrix

Idea: Instead of arranging splitters in 2^n binary tree, arrange them in *$n \times n$ matrix*:



5x5 splitter matrix


Space complexity n^2 . Step complexity?

Matrix Collect

Step complexity of first *STORE* is $O(k)$, subsequent ones are $O(1)$. *COLLECT* has step complexity $O(k^2)$.

Proof.

- Let x_i be number of procs entering row i . By induction on i , $x_i \leq k-i$.
- Of course: $x_0 \leq k$
- Let j be largest column s.t. at least one process visits the splitter at $(i-1, j)$.
- Not all processes go left, so $x_i \leq k-i$.
- Same for column.
- So *every process stops the latest in row $k-1$ and column $k-1$.*
- The number of marked splitters is at most k^2 : complexity of COLLECT.
- The longest path in matrix is $2k$, so STORE complexity at most $O(k)$.



QED

Remarks

- Randomized algorithms can achieve binary trees of depth $O(\log n)$ only.
- $O(k)$ step complexity and $O(n^2)$ space complexity is possible for COLLECT, even deterministically

End of Lecture