



Consensus




Consensus




Assumptions in this lecture

- All-to-all communication




Assumptions in this lecture

- All-to-all communication
- Two possible input values, 0 and 1




Assumptions in this lecture

- All-to-all communication
- Two possible input values, 0 and 1
- Unique IDs




Assumptions in this lecture

- All-to-all communication
- Two possible input values, 0 and 1
- Unique IDs
- Asynchronous model




Recall: Asynchronous model


- Messages arrive eventually
we cannot assume an upper bound on the message delay



Consensus




Consensus



Crash failures


- A node can crash at any time,
- This node does not recover anymore,
- Messages that have been successfully sent by this node arrive eventually,
- We know how many nodes can crash – f of them



Requirements


- *Agreement*
the nodes agree on the same value
- *Termination*
the nodes terminate in a finite time
- *Validity*
the decision should be one of the inputs

Examples



**No deterministic algorithm
in the asynchronous model
can solve consensus**


Algorithm Configuration



Fully described by:

- States
- Messages in transit


Initial Configuration C_0



Fully described by:

- IDs
- Inputs of the nodes


Bivalent Configuration B_i



Configuration after which **both decisions**, 0 and 1, can follow


Univalent Configurations

U_0 and U_1




Configuration after which only
one decision can follow, e.g., 1

Critical Configuration




Last configuration after which
both decisions can follow

Transition τ




A transition $\tau = (m, u)$ is characterized by a node u receiving a message m


Configuration tree



Configuration tree




Configuration tree




bivalent configuration

Configuration tree



Configuration tree




Impossibility result - recipe

- There always exists a bivalent initial configuration
- There must exist a critical configuration
- The action of a single node decides whether the outcome is 0/1


Impossibility result - recipe

- There always exists a bivalent initial configuration
- There must exist a critical configuration
- The action of a single node decides whether the outcome is 0/1

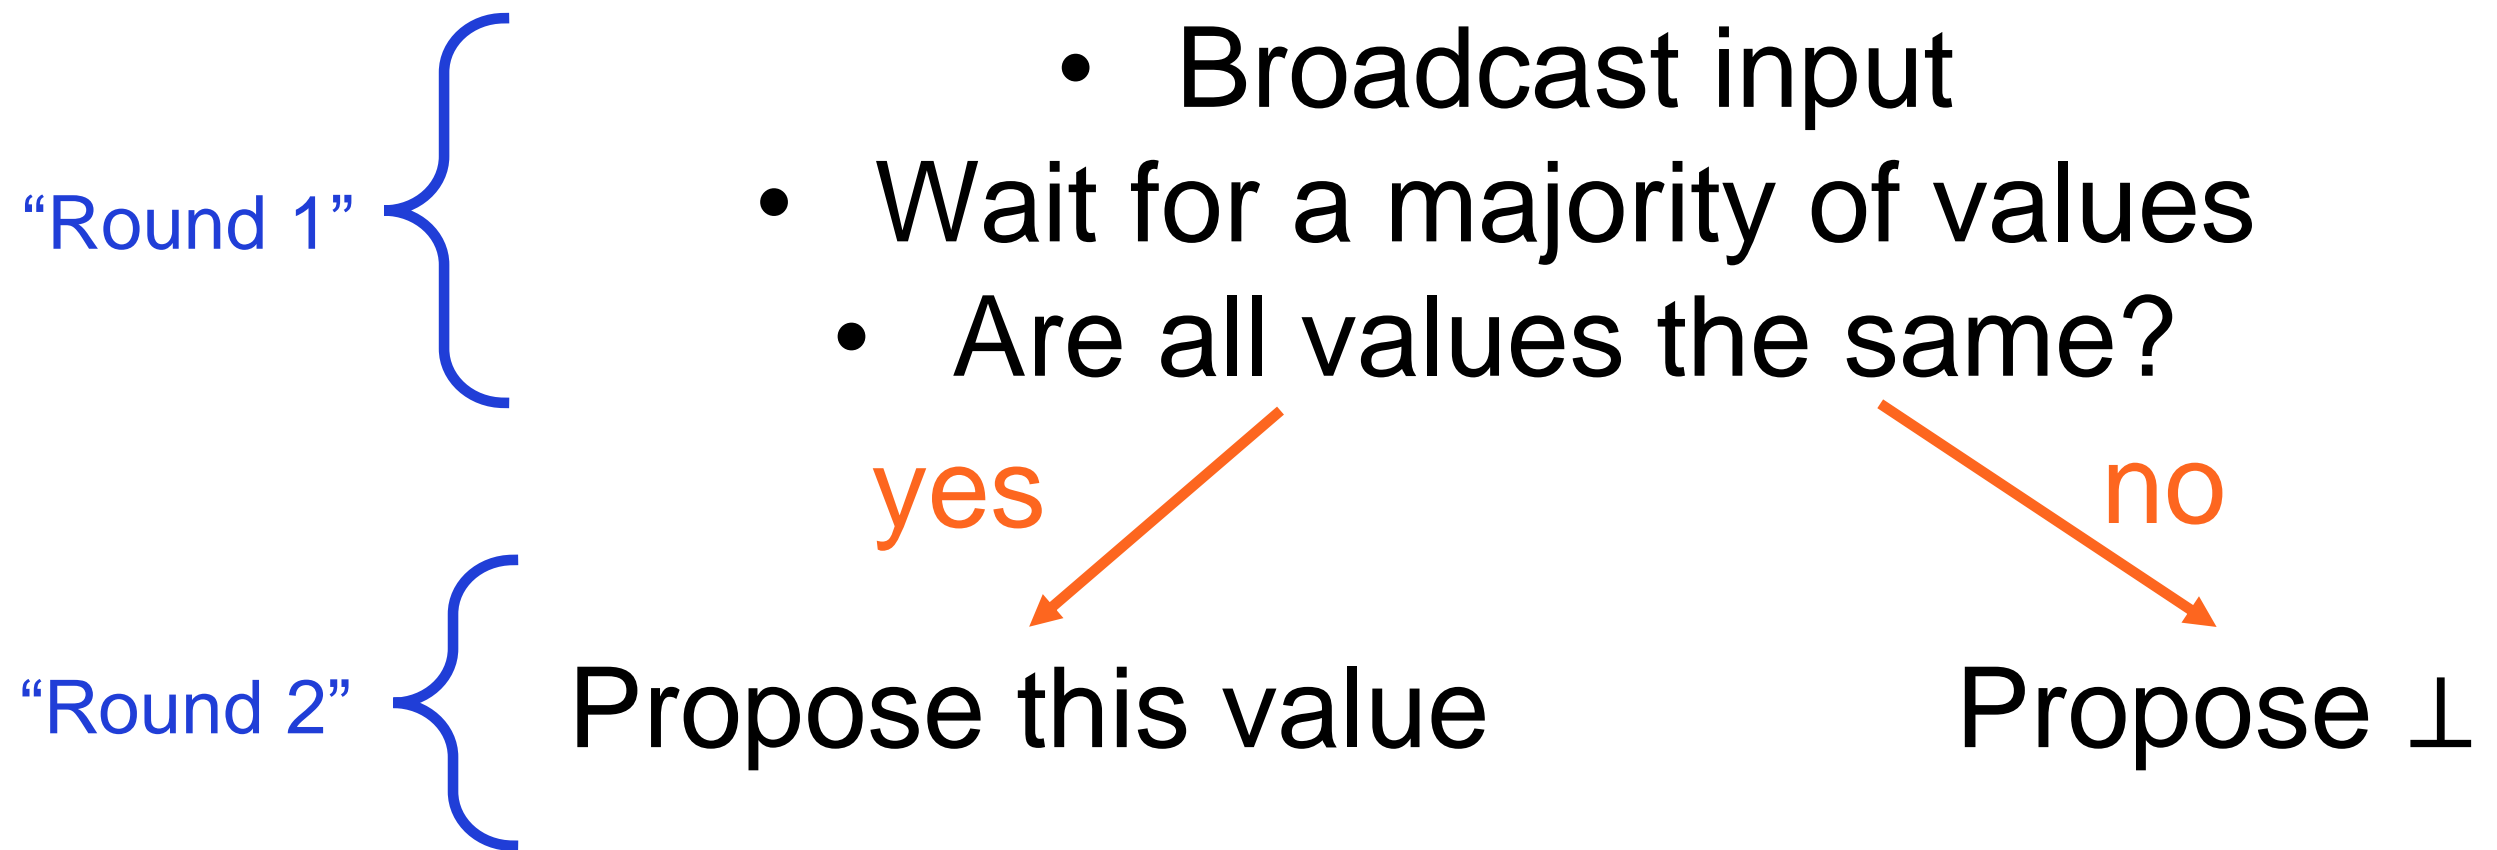


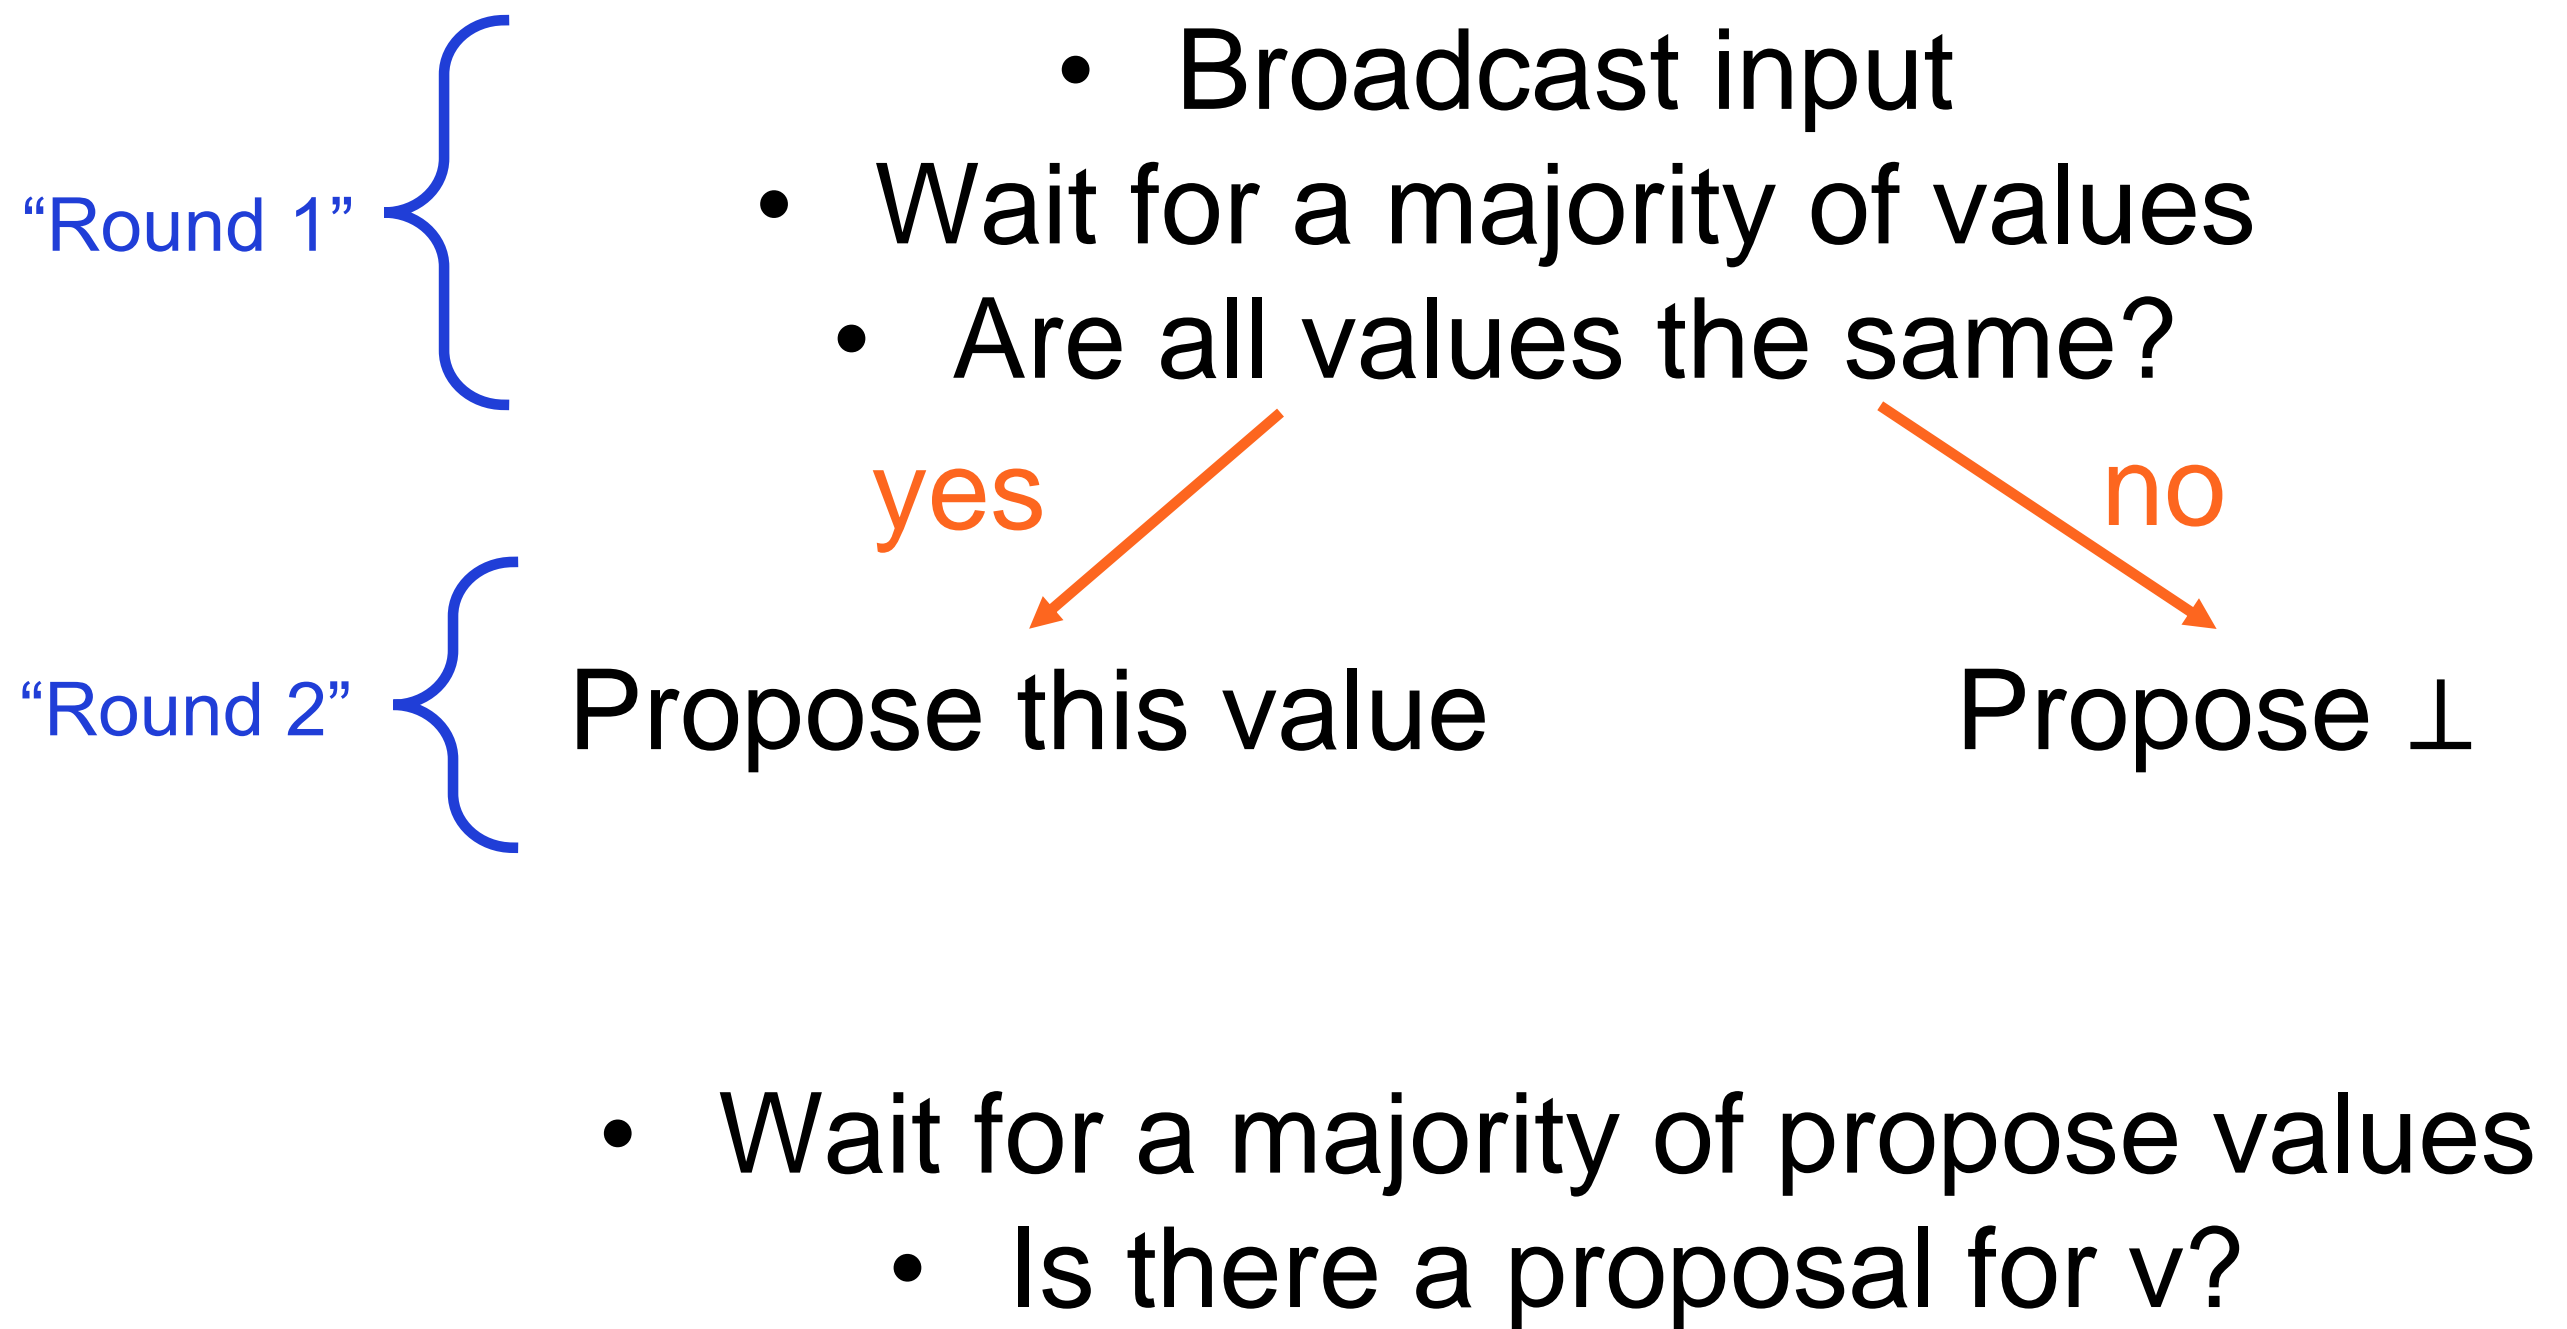
Randomized Consensus, $f < \frac{n}{2}$

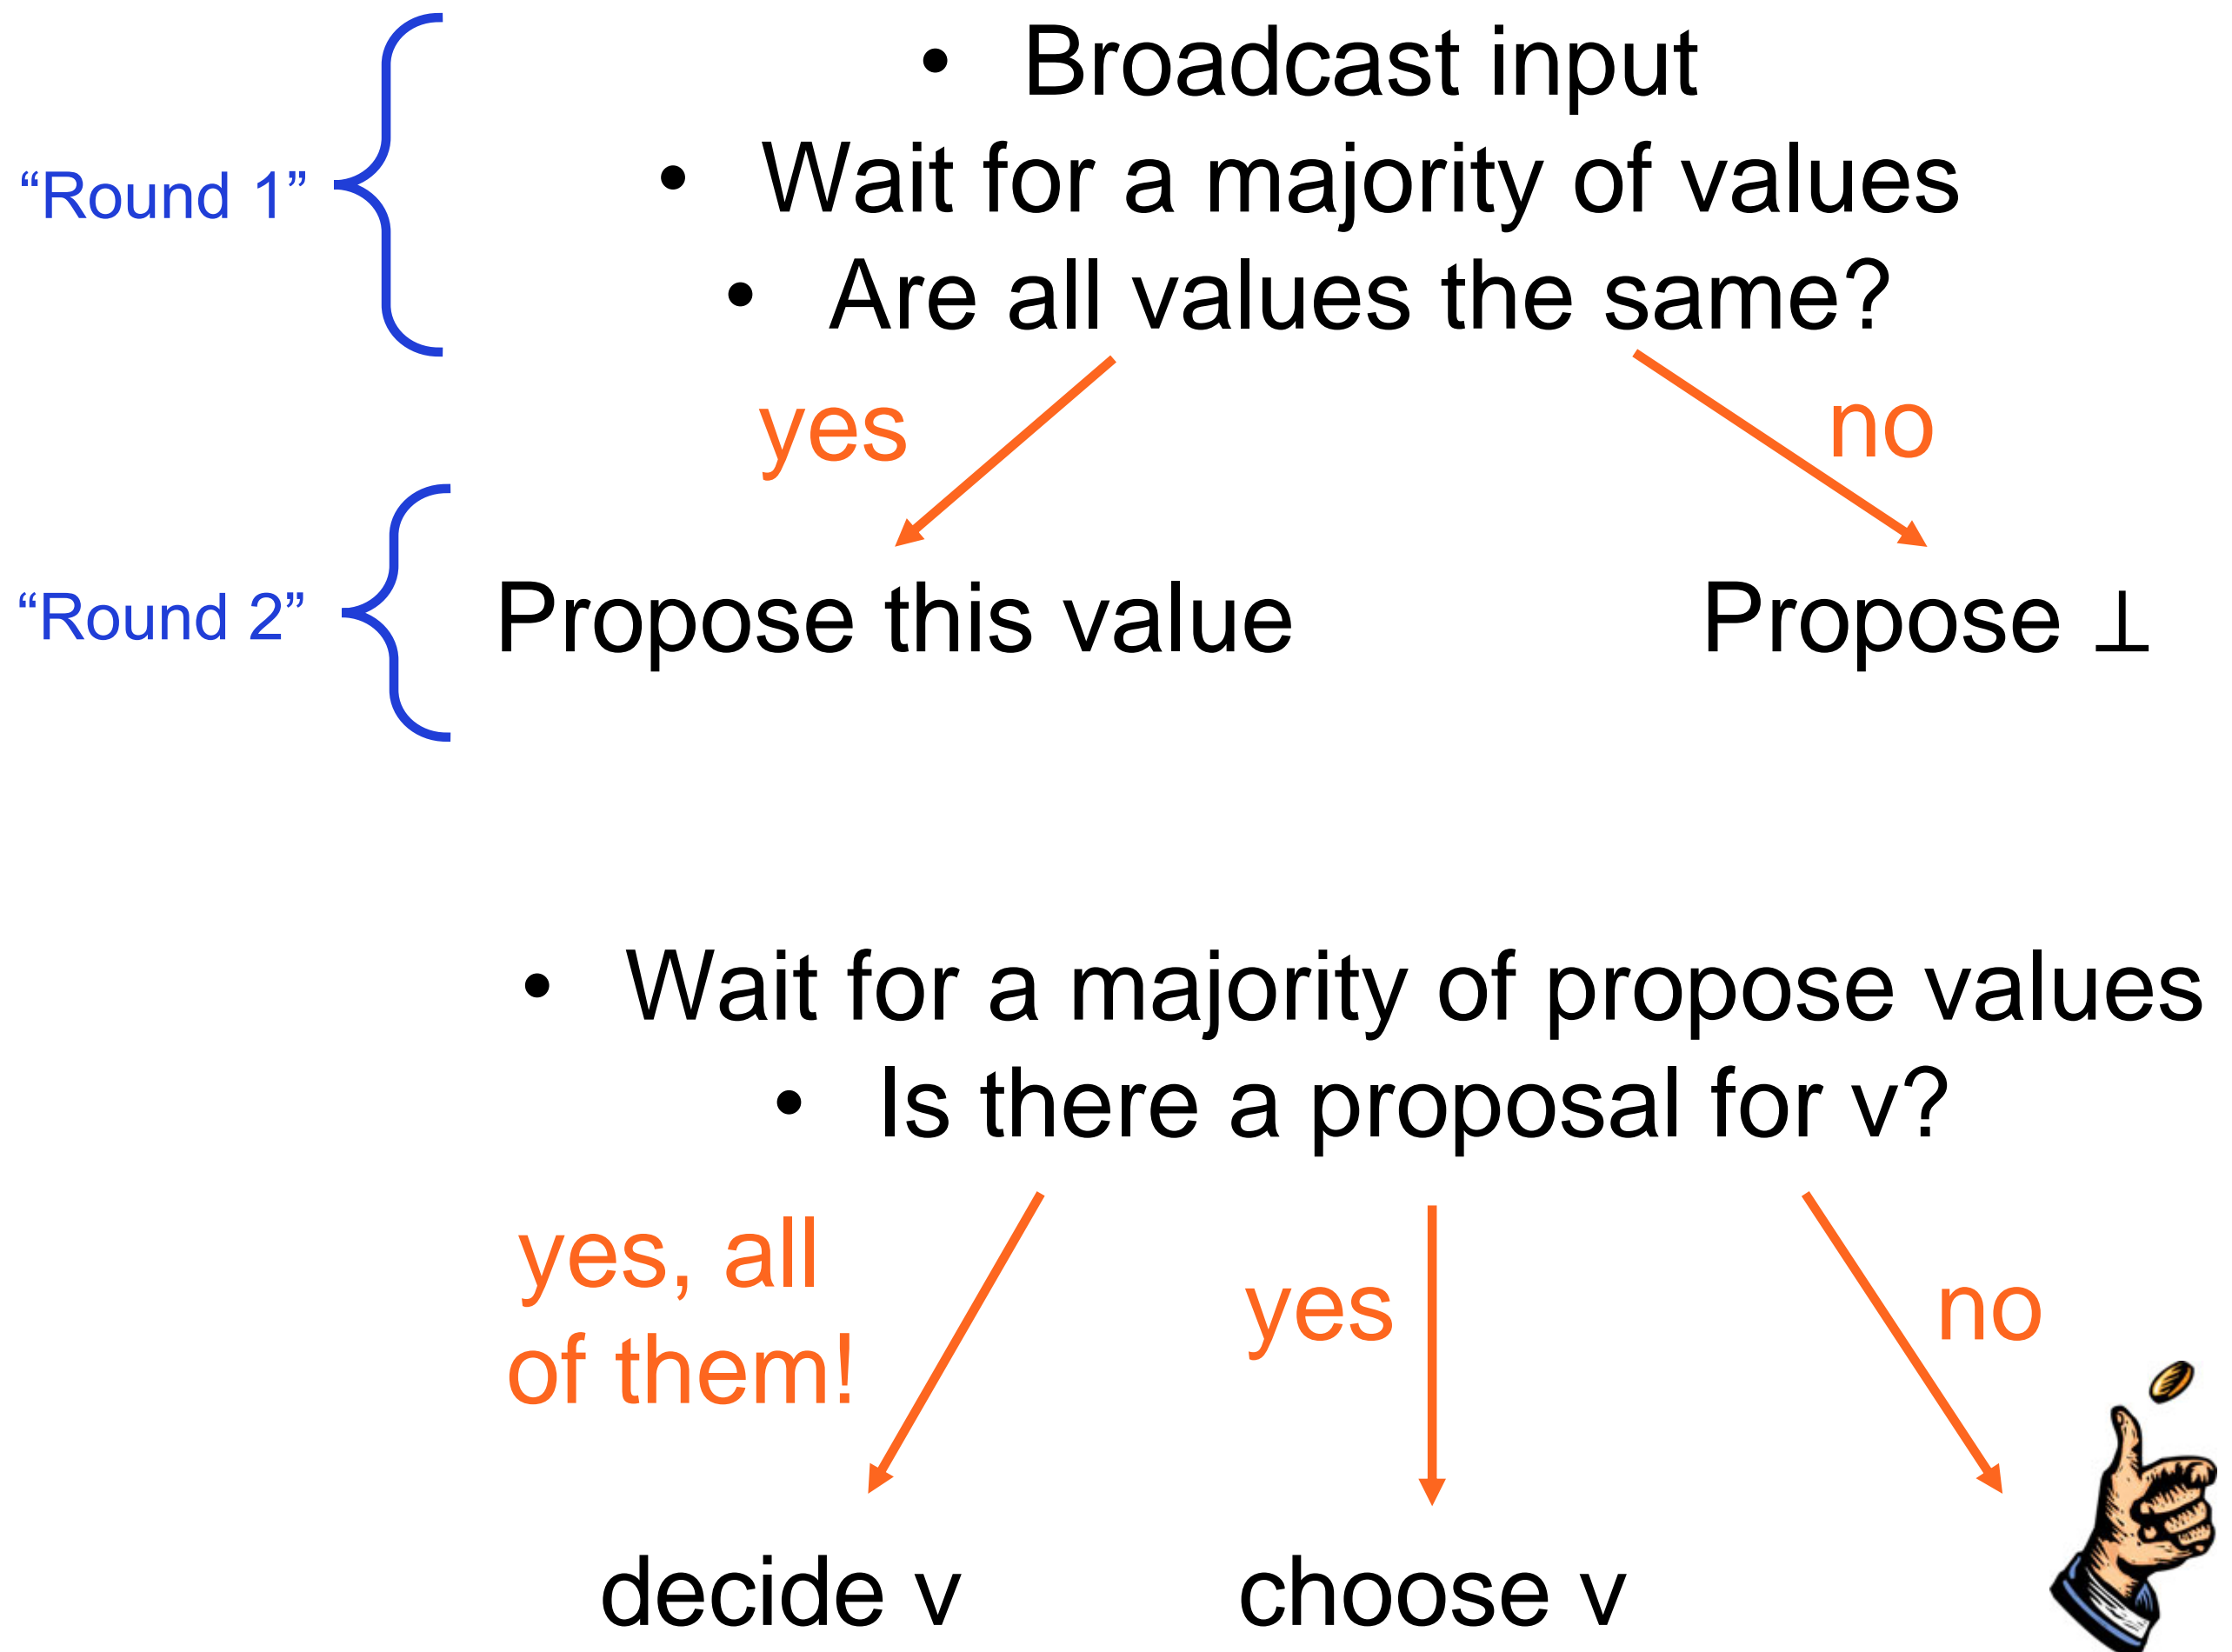
“Round 1”




- Broadcast input
- Wait for a majority of values
- Are all values the same?










Example



Can any algorithm handle $f = \frac{n}{2}$ crashes?

Example



Improving the coinflip, $f < \frac{n}{3}$

Requirement:
output 0 and 1 with constant probability

- Biased local coin:
 - 0 with probability $1/n$
 - 1 with probability $1 - \frac{1}{n}$
- Broadcast coinflip


- Biased local coin:
 - 0 with probability $1/n$
 - 1 with probability $1 - \frac{1}{n}$
- Broadcast coinflip
- Wait for $n - f$ coins, store them in C_u
- Broadcast C_u

- Biased local coin:
 - 0 with probability $1/n$
 - 1 with probability $1 - \frac{1}{n}$
- Broadcast coinflip
- Wait for $n - f$ coins, store them in C_u
- Broadcast C_u
- Wait for $n - f$ coin sets
- Does one of them contain 0?

- Biased local coin:
 - 0 with probability $1/n$
 - 1 with probability $1 - \frac{1}{n}$
- Broadcast coinflip
- Wait for $n - f$ coins, store them in C_u
- Broadcast C_u
- Wait for $n - f$ coin sets
- Does one of them contain 0?



Analysis



Learning goals

- **Problems:** Consensus
- **Distributed models:** asynchronous all-to-all communication
- **Impossibility results:**
 - Impossibility of deterministic asynchronous consensus
 - Impossibility of consensus with $n/2$ failures
- **Algorithms:**
 - Randomized consensus algorithm with $f < n/2$
 - Biased local coin